

# Graph Neural Networks

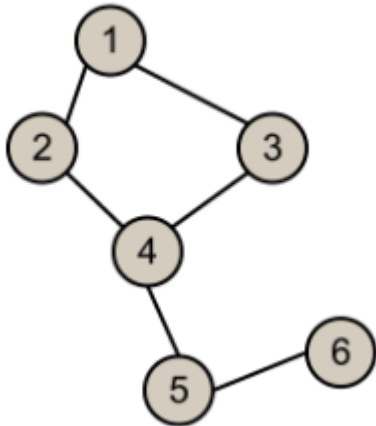
# Outline

- Introduction to graphs and where to find them
- Machine learning tasks on graphs
- Introduction to graph neural networks (GNNs)
- The computational graph of a GNN
- Convolutional Graph Neural Networks
- Introduction to Pytorch Geometric
- Node Classification with Pytorch Geometric
- Graph Attention Networks.
- The over smoothing problem
- Graph Representation learning
- Graph Classification with Pytorch Geometric

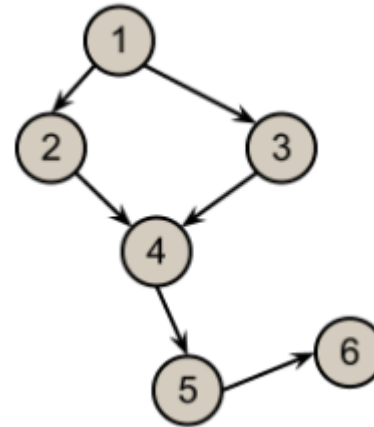
# Graphs

A **graph** is an ordered pair  $(V,E)$  where,

- $V$  is the *vertex set* (also *node set* ) whose elements are the vertices, or *nodes* of the graph.
- $E$  is the *edge set* whose elements are the edges, or connections between vertices, of the graph. If the graph is undirected, individual edges are unordered pairs.
- If the graph is directed, edges are ordered pairs

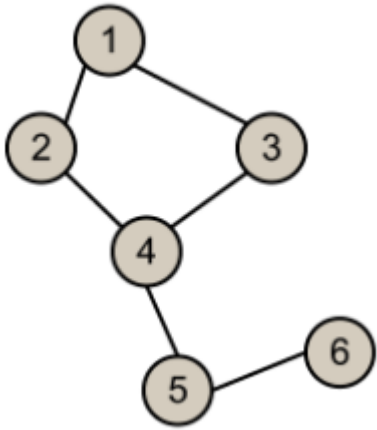


undirected



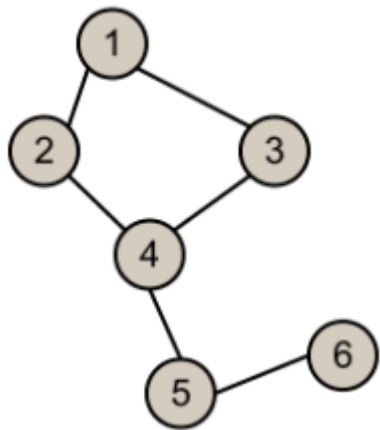
directed

# Graphs representation: adjacency matrix

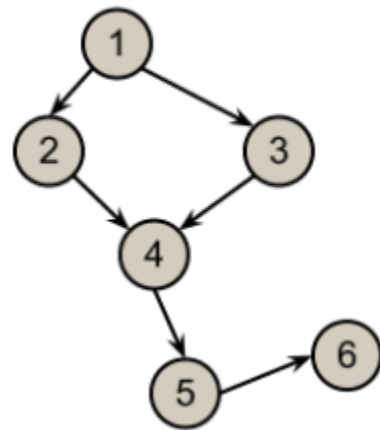


	1	2	3	4	5	6
1	0	1	1	0	0	0
2	1	0	0	1	0	0
3	1	0	0	1	0	0
4	0	1	1	0	1	0
5	0	0	0	1	0	1
6	0	0	0	0	1	0

# Graphs representation: adjacency matrix



	1	2	3	4	5	6
1	0	1	1	0	0	0
2	1	0	0	1	0	0
3	1	0	0	1	0	0
4	0	1	1	0	1	0
5	0	0	0	1	0	1
6	0	0	0	0	1	0

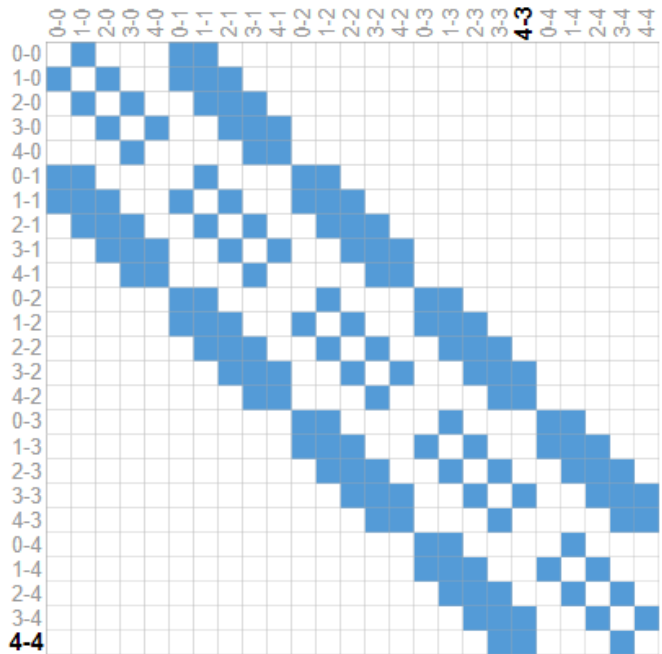


	1	2	3	4	5	6
1	0	1	1	0	0	0
2	-1	0	0	1	0	0
3	-1	0	0	1	0	0
4	0	-1	-1	0	1	0
5	0	0	0	-1	0	1
6	0	0	0	0	-1	0

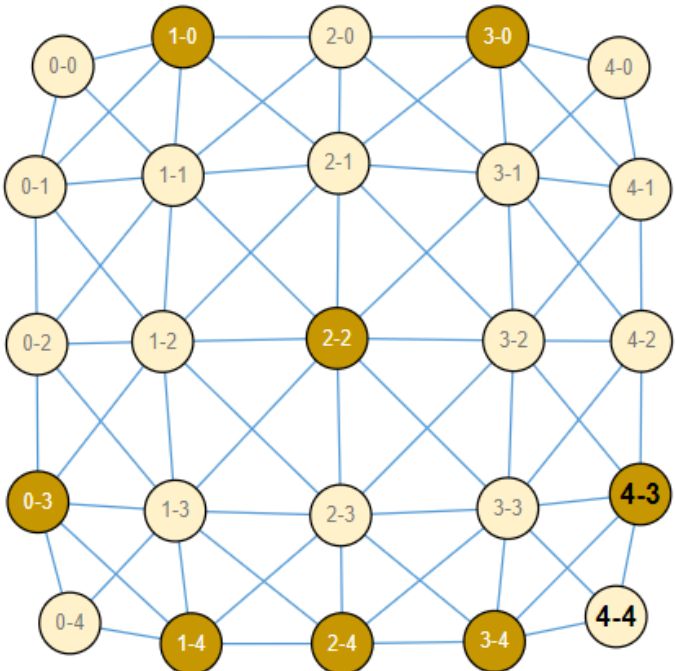
# Graphs and where to find them

0-0	1-0	2-0	3-0	4-0
0-1	1-1	2-1	3-1	4-1
0-2	1-2	2-2	3-2	4-2
0-3	1-3	2-3	3-3	4-3
0-4	1-4	2-4	3-4	4-4

Image Pixels



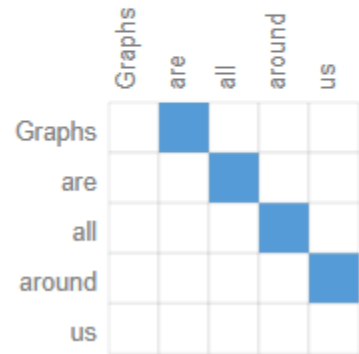
Adjacency Matrix



Graph

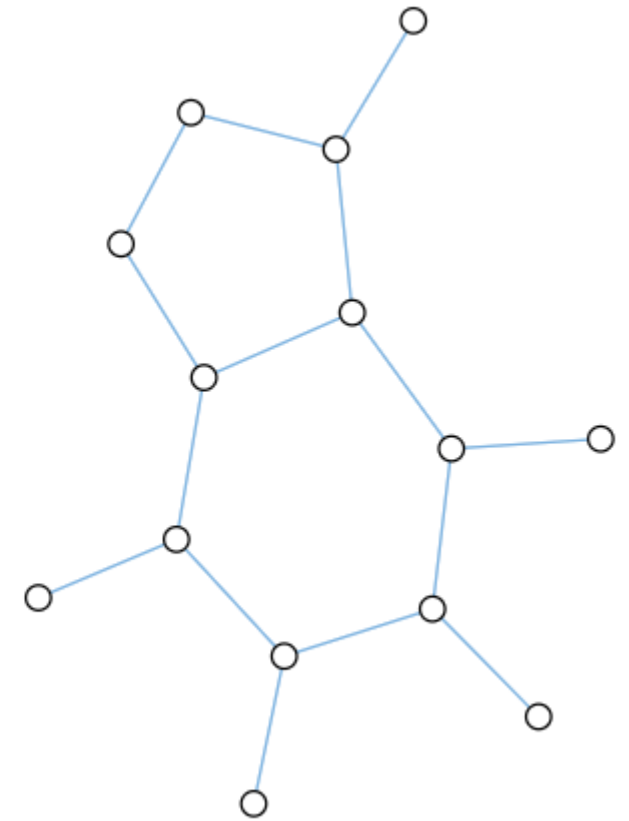
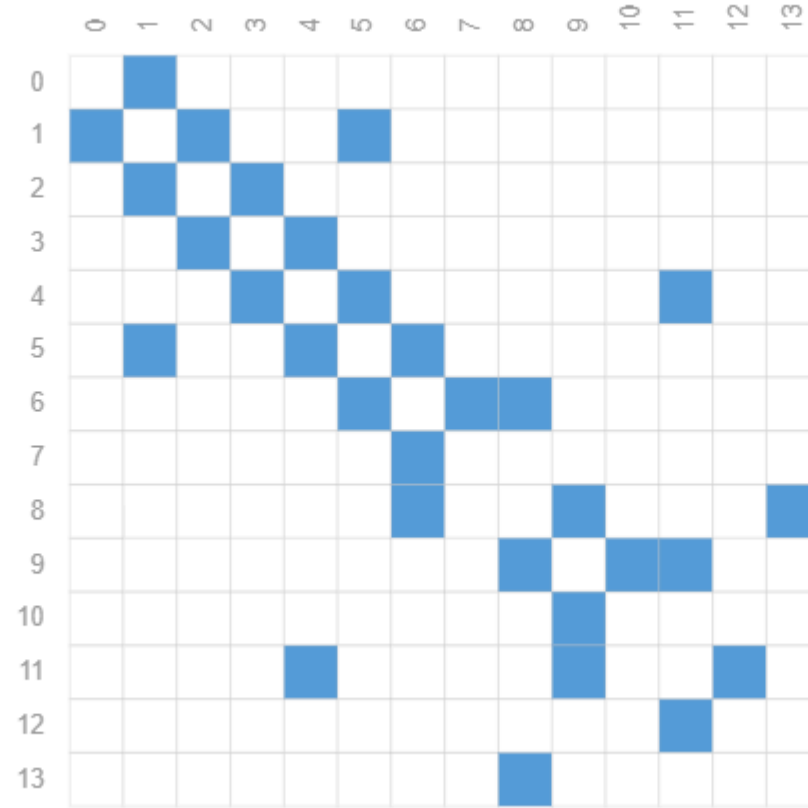
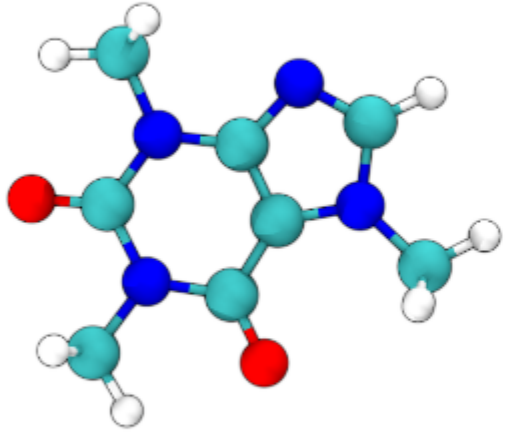
## Images as graphs

Graphs and where to find them



Text as graphs

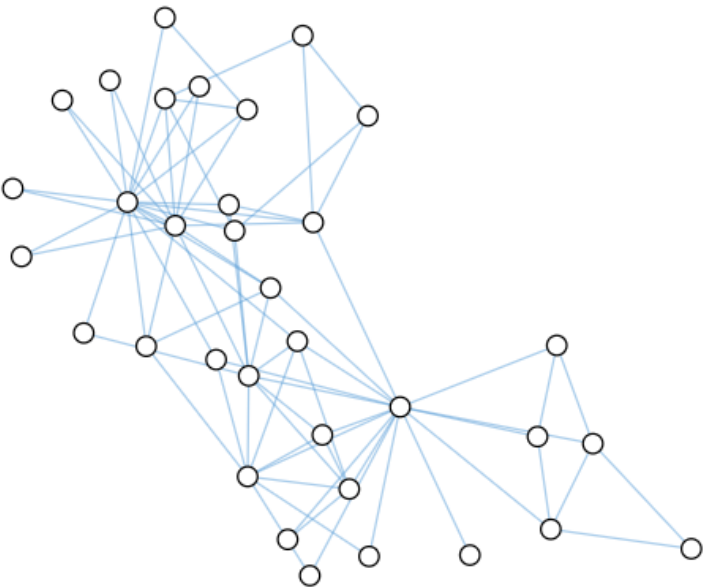
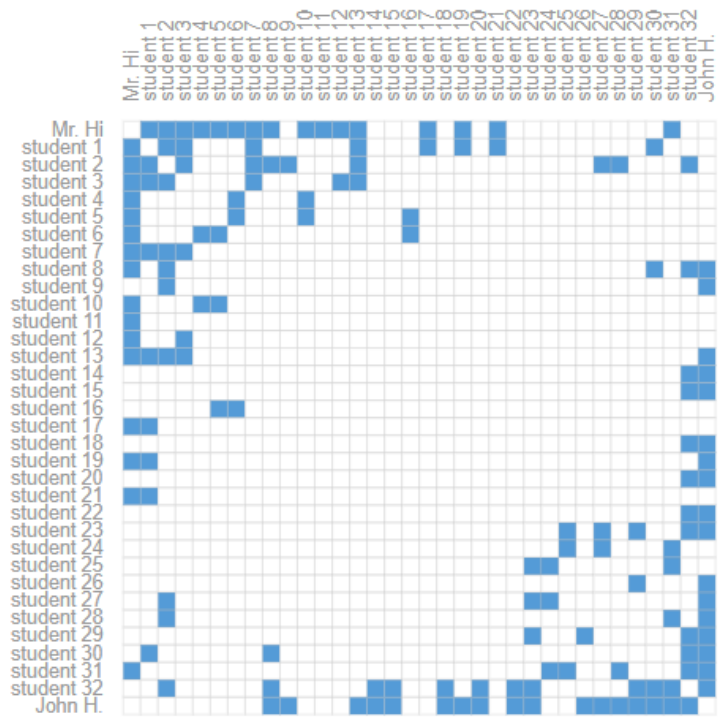
## Graphs and where to find them



(Left) 3d representation of the Caffeine molecule (Center) Adjacency matrix of the bonds in the molecule (Right) Graph representation of the molecule.

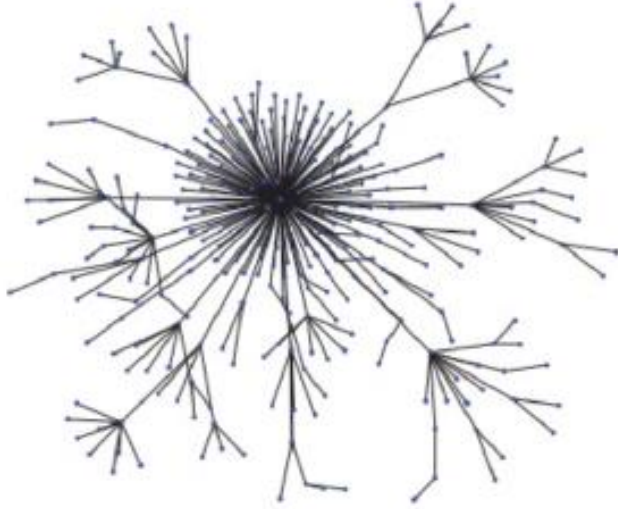


# Other examples

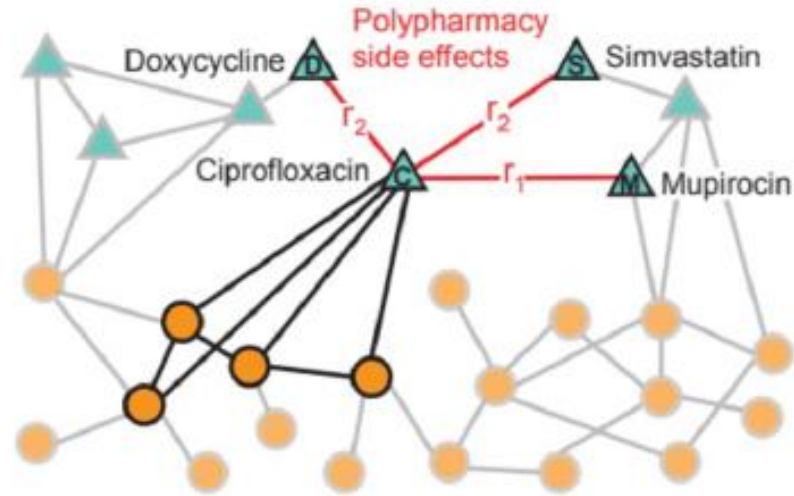


Social networks as graphs

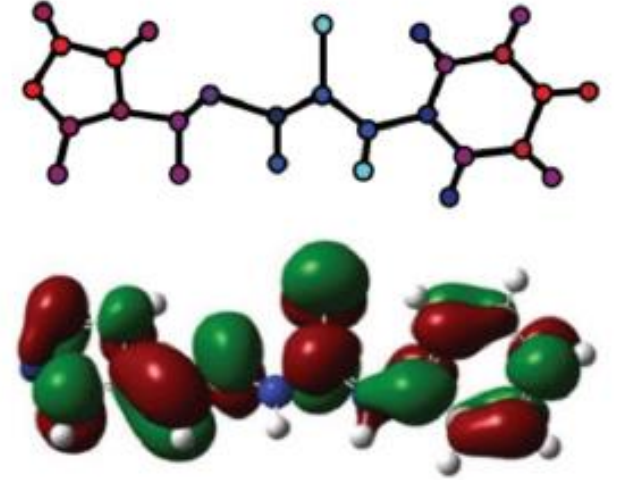
## Other examples



Fake news detection



Polypharmacy



Chemistry

## Other examples



Social Graphs



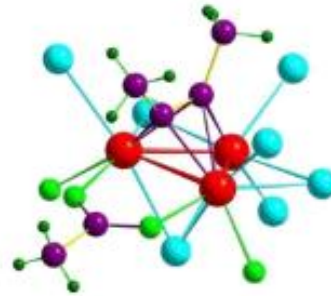
Transportation Graphs



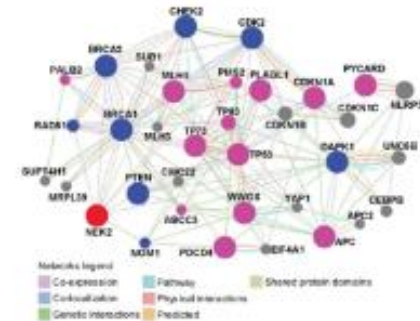
Brain Graphs



Web Graphs



Molecular Graphs

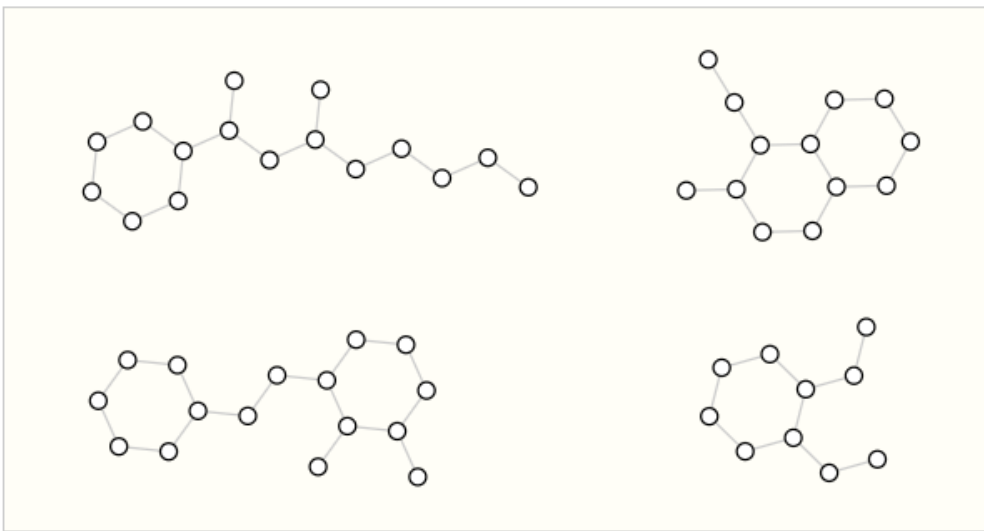


Gene Graphs

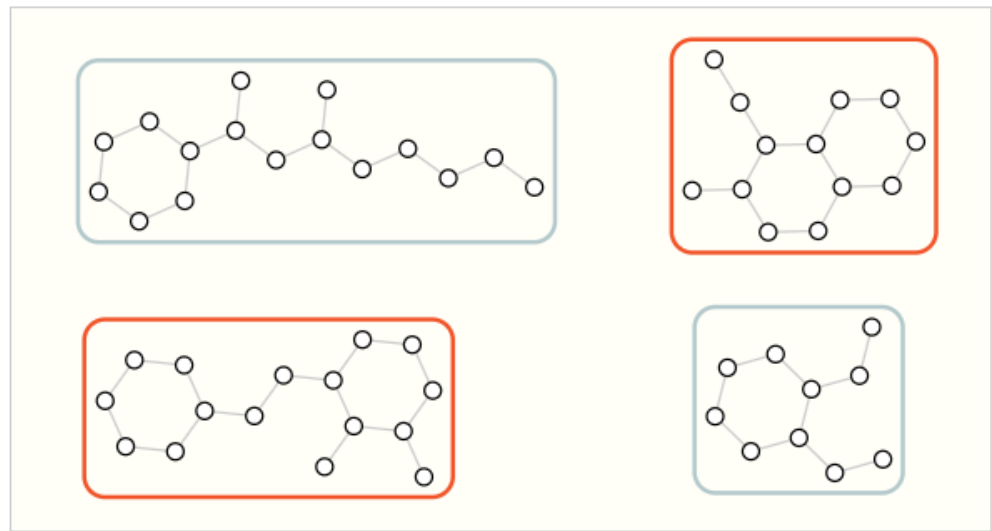
Protein-protein interactions , Citation networks as graphs, Machine learning models, etc.

## ML on graphs: Graph level task

In a graph-level task, our goal is to predict the property of an entire graph.



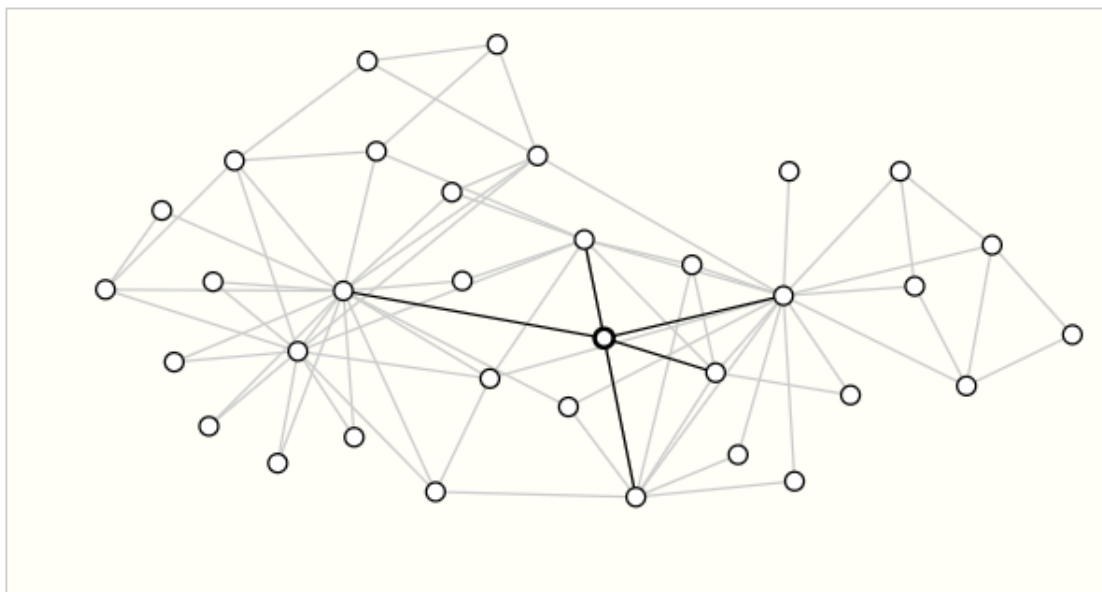
**Input:** graphs



**Output:** labels for each graph, (e.g., "does the graph contain two rings?")

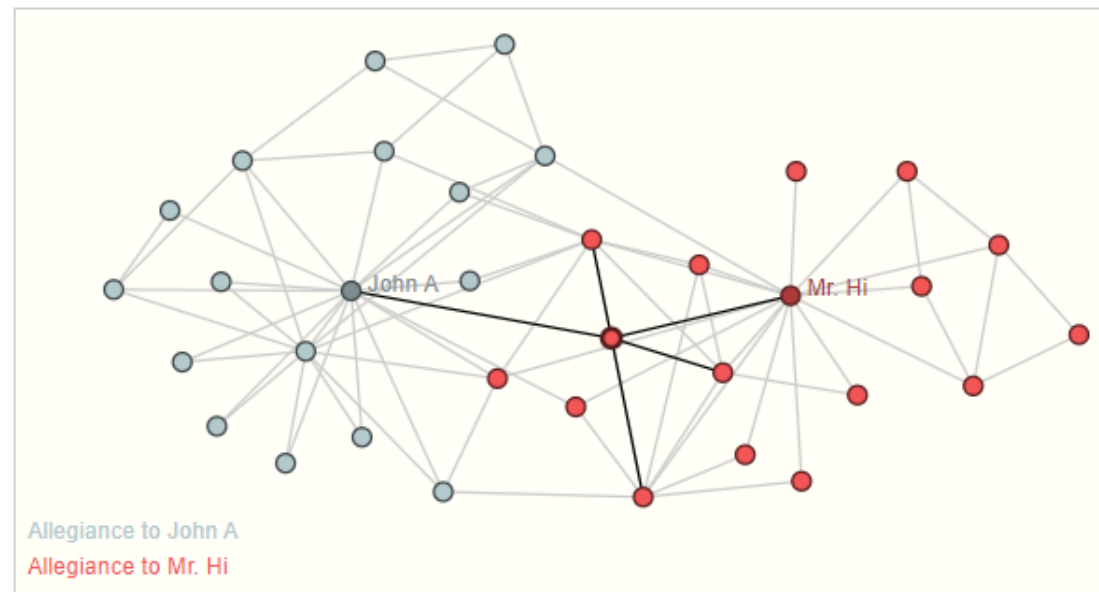
## ML on graphs: Node-level task

In a graph-level task, our goal is to predict the property of an entire graph.



**Input:** graph with unlabeled nodes

→

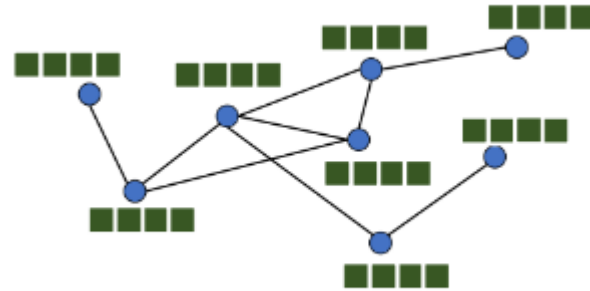


**Output:** graph node labels

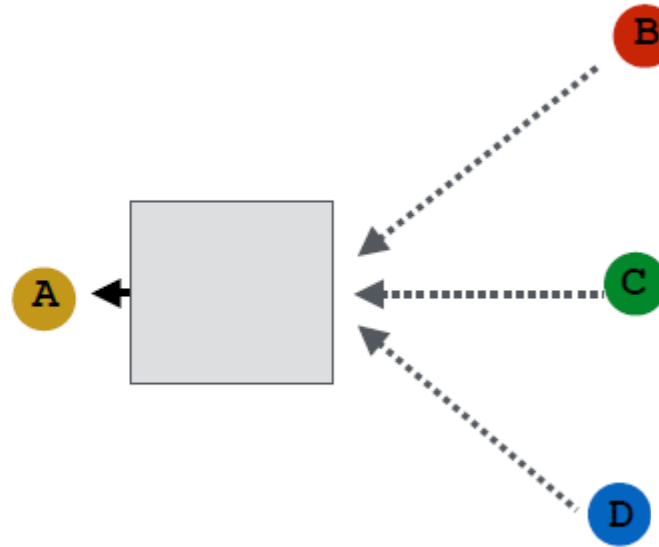
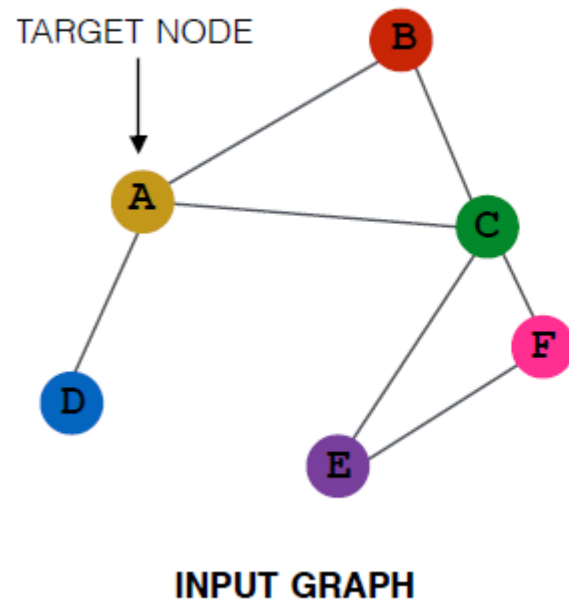
## Problem setup

Assume we have a graph  $G$ :

- $V$  is the **vertex set**
- $A$  is the **adjacency matrix** (assume binary)
- $X \in \mathbb{R}^{m \times |V|}$  is a matrix of **node features**
- $v$ : a node in  $V$ ;  $N(v)$ : the set of neighbors of  $v$ .



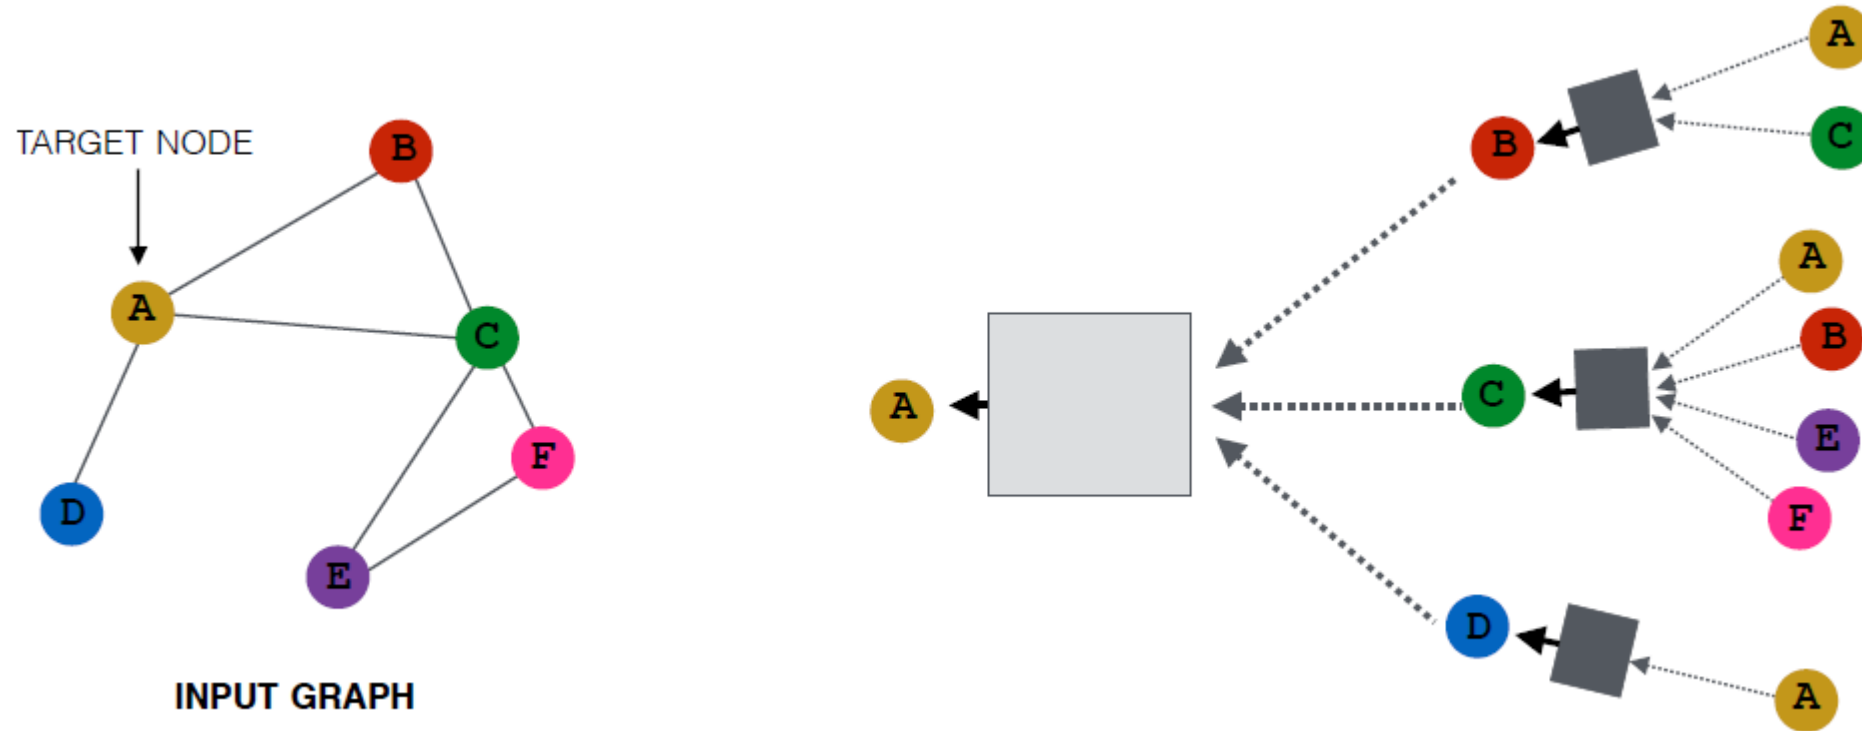
## Graph Neural Networks



Key idea: Each node defines a computation graph

- Each edge in this graph is a transformation/aggregation function

## Graph Neural Networks



Key idea: Each node defines a computation graph

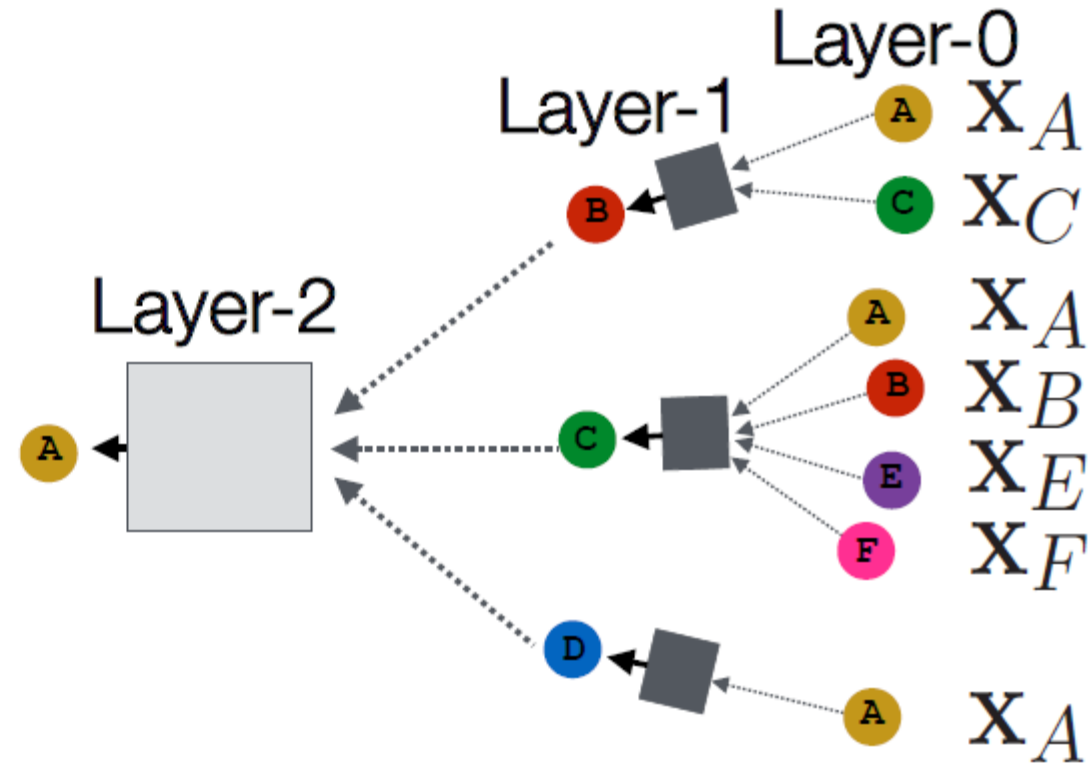
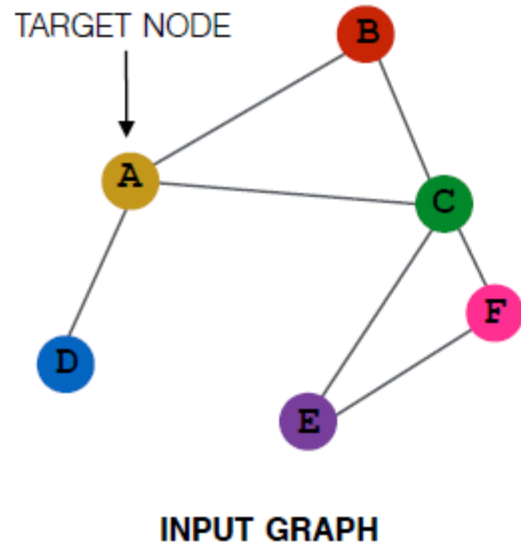
- Each edge in this graph is a transformation/aggregation function



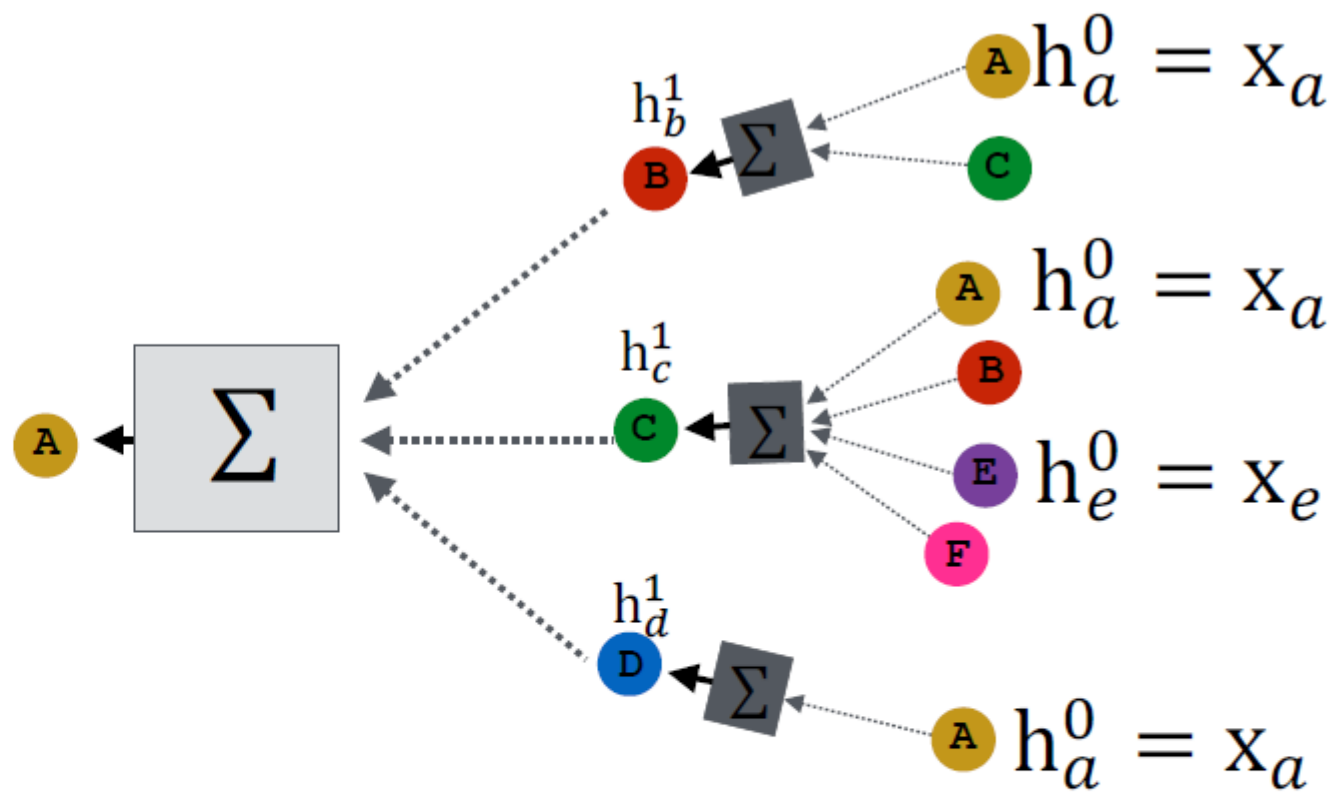
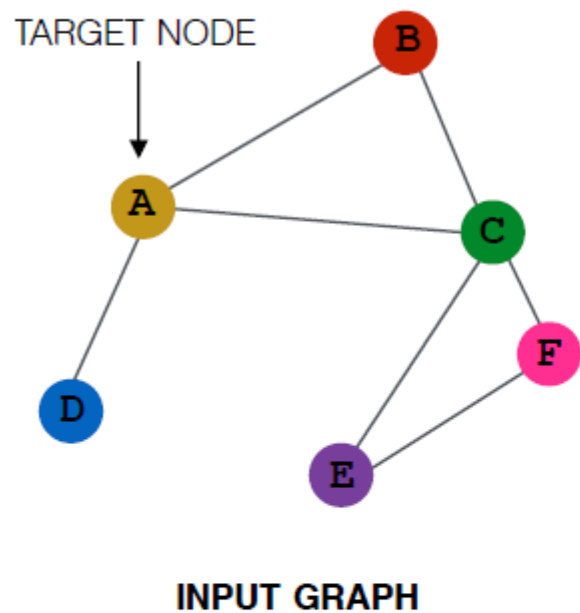
## Graph Neural Networks

Nodes have embeddings at each layer

- Layer-0 embedding of node  $u$  is its input feature,  $x_u$
- Layer- $k$  embedding gets information from nodes that are  $k$  hops away

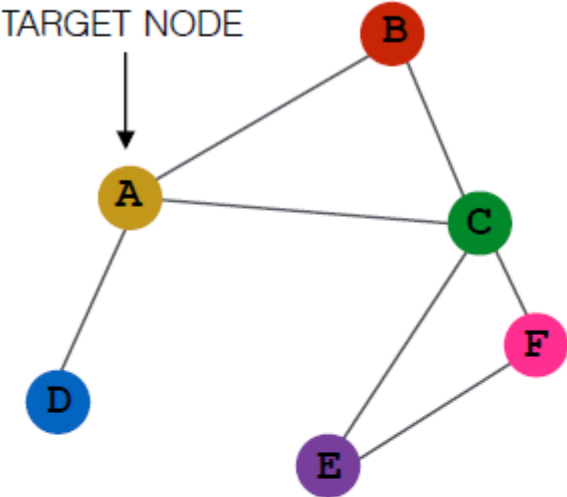


## Graph Neural Networks

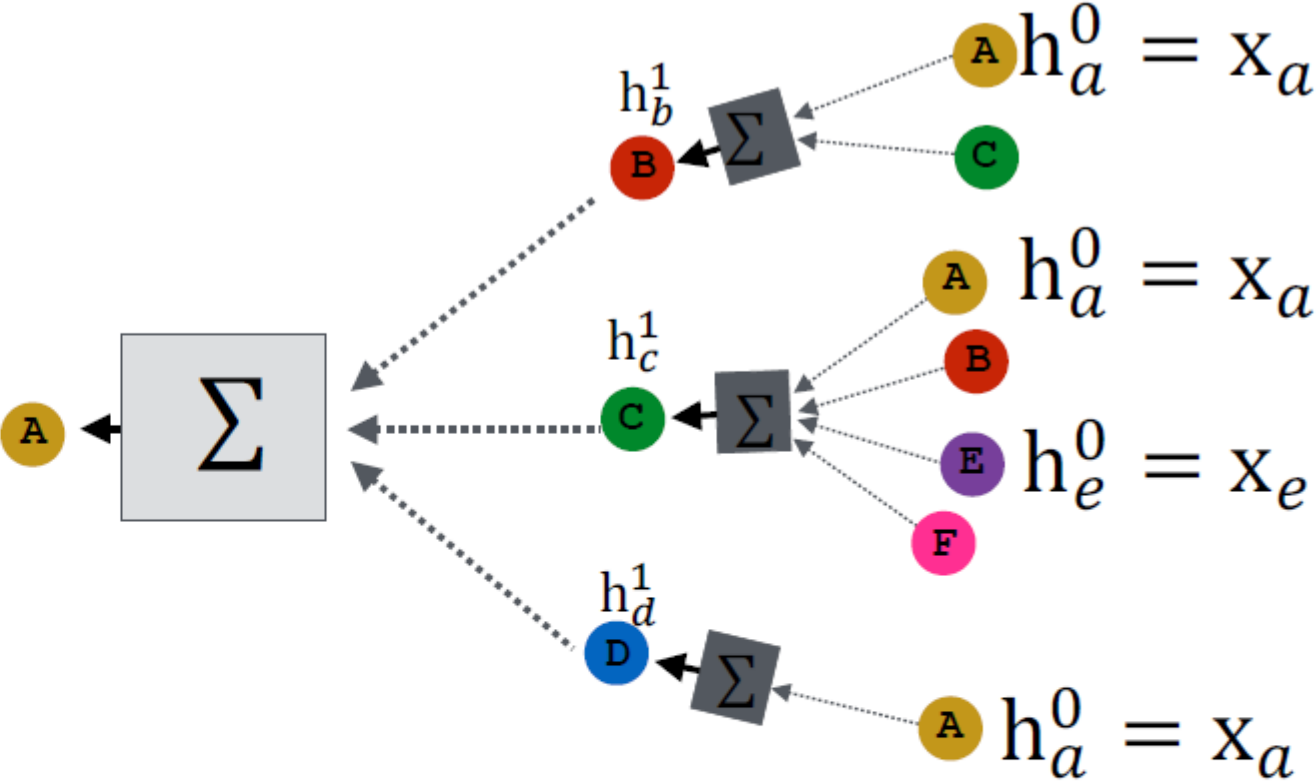


$$h_v^{(l+1)} = \sigma(W_l \sum_{u \in N(v)} \frac{h_u^{(l)}}{|N(v)|} + B_l h_v^{(l)})$$

# Graph Neural Networks



INPUT GRAPH

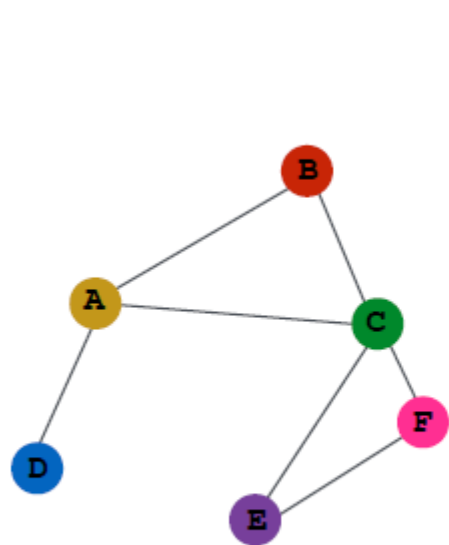


$$h_v^{(l+1)} = \sigma(W_l \sum_{u \in N(v)} \frac{h_u^{(l)}}{|N(v)|} + B_l h_v^{(l)})$$

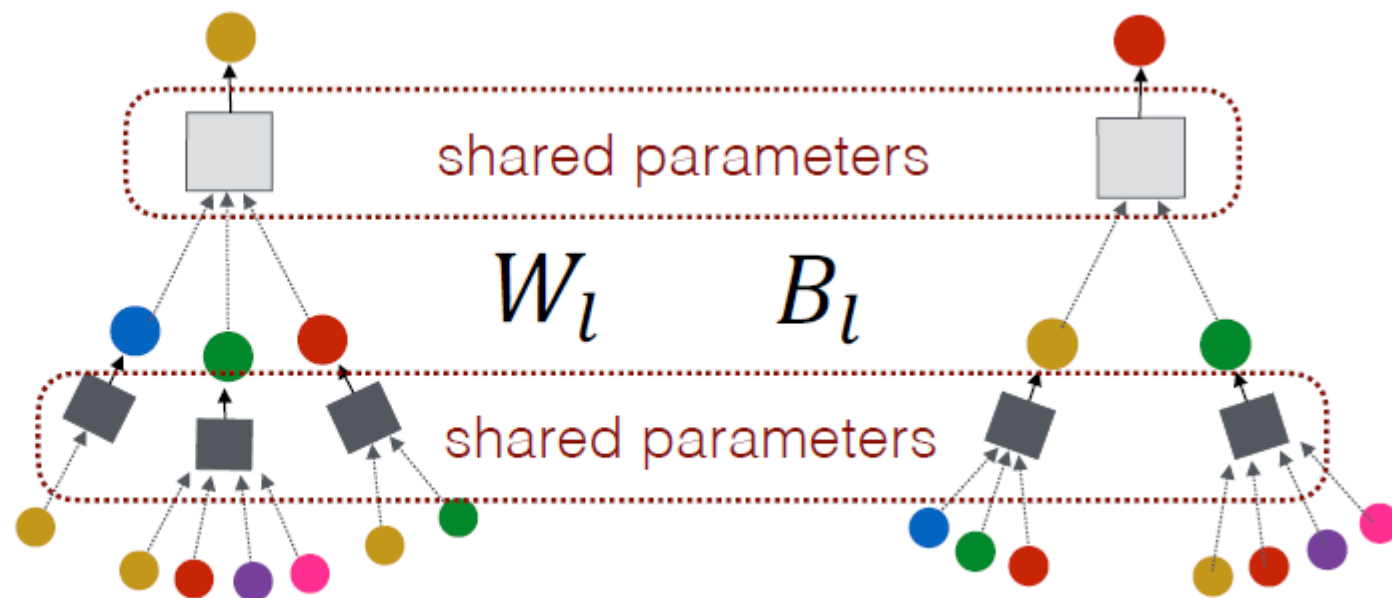
Learnable parameters

## Graph Neural Networks

The *same* aggregation parameters are shared for all nodes



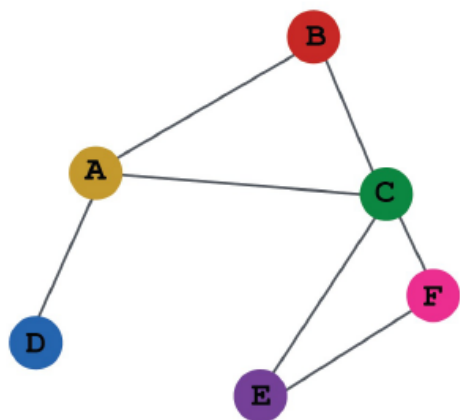
INPUT GRAPH



Compute graph for node A

Compute graph for node B

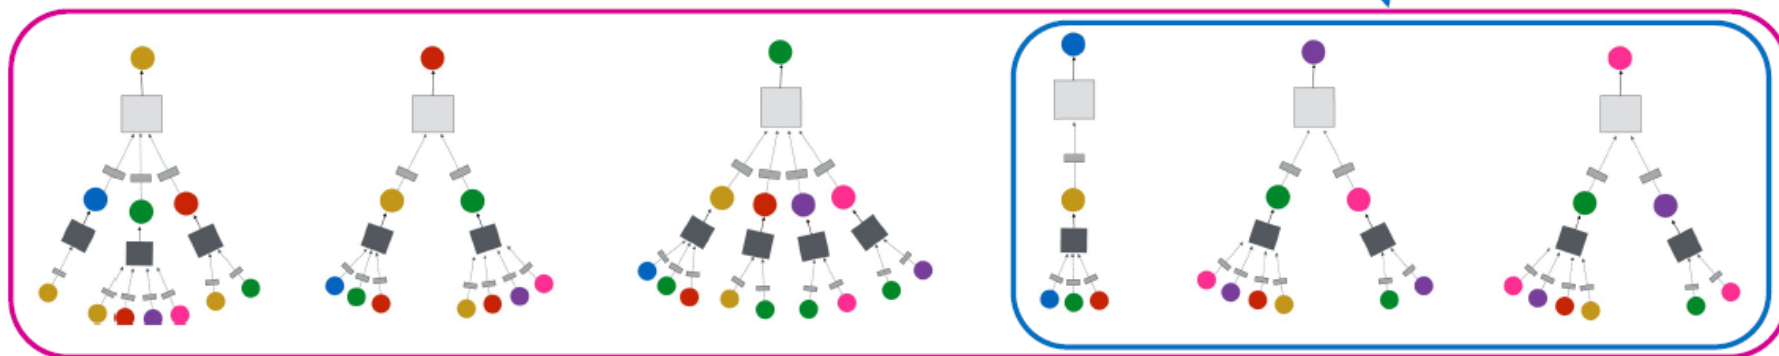
# Graph Neural Networks



INPUT GRAPH

Generate embeddings  
for nodes as needed

Even for nodes we never  
trained on!



## Key Benefits

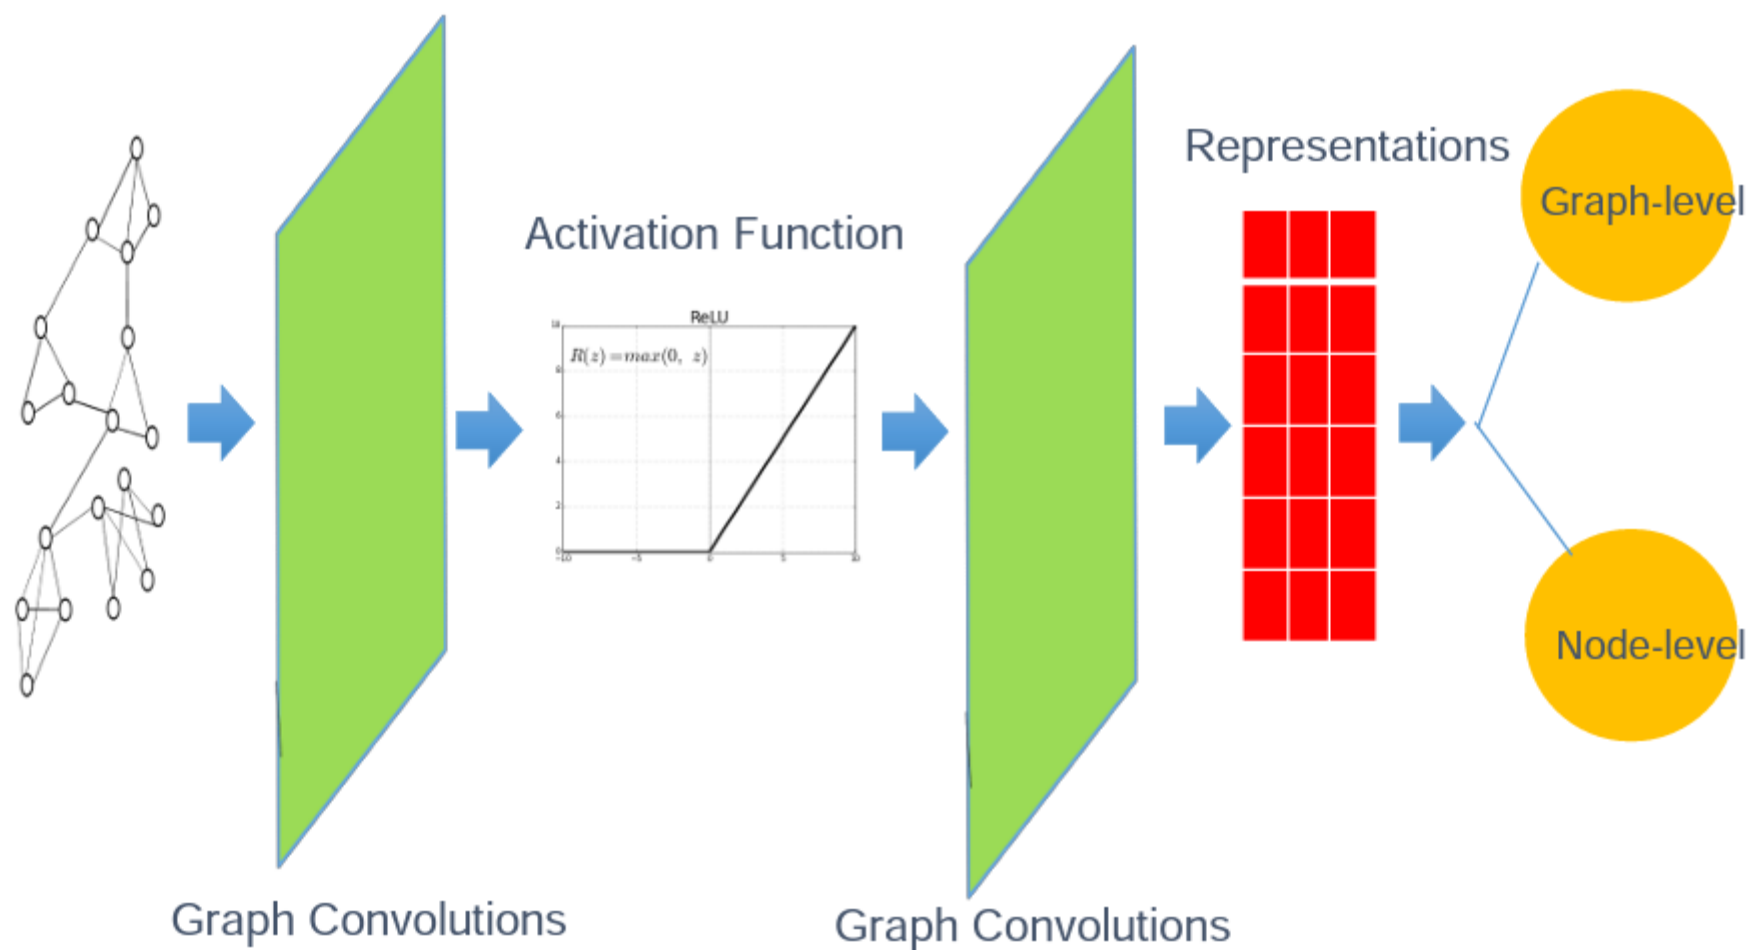
- \*No manual feature engineering needed
- \*End-to-end learning results in optimal features.
- \*Any graph machine learning task: Node-level, link-level, entire graph-level prediction
- \*Scalable to billion node graphs!

## Graph Convolutional Network

$$h_v^{k+1} = \sigma(W_k \sum_{u \in N(v)} \frac{h_u^k}{|N(v)|} + B_k h_v^k)$$

**Key benefit : fast and easy to implement**

# Graph Neural Networks





## Pytorch Geometric

PyG (PyTorch Geometric) is a library built upon PyTorch to easily write and train Graph Neural Networks (GNNs) for a wide range of applications related to structured data.



[https://pytorch-geometric.readthedocs.io/en/latest/get\\_started/colabs.html](https://pytorch-geometric.readthedocs.io/en/latest/get_started/colabs.html)

## Pytorch Geometric

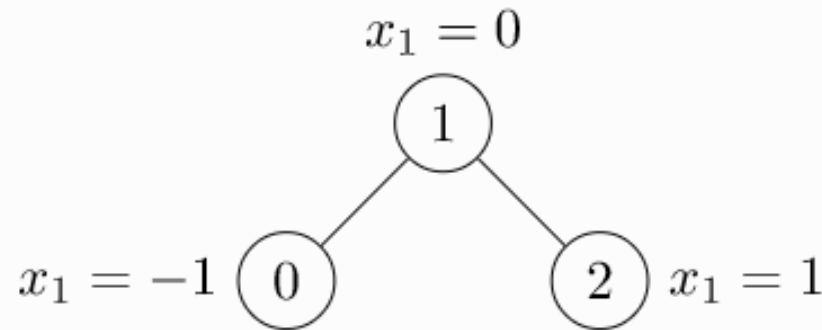
```
import torch
from torch_geometric.data import Data

edge_index = torch.tensor([[0, 1, 1, 2],
                           [1, 0, 2, 1]], dtype=torch.long)
x = torch.tensor([[-1], [0], [1]], dtype=torch.float)

data = Data(x=x, edge_index=edge_index)
>>> Data(edge_index=[2, 4], x=[3, 1])
```

# Graph data structure in pytorch Geometric

# Input data on the graph




Although the graph has only two edges, we need to define four index tuples to account for both directions of an edge

[https://pytorch-geometric.readthedocs.io/en/latest/get\\_started/introduction.html](https://pytorch-geometric.readthedocs.io/en/latest/get_started/introduction.html)

## Pytorch Geometric

```
import torch
from torch.nn import Linear
from torch_geometric.nn import GCNConv

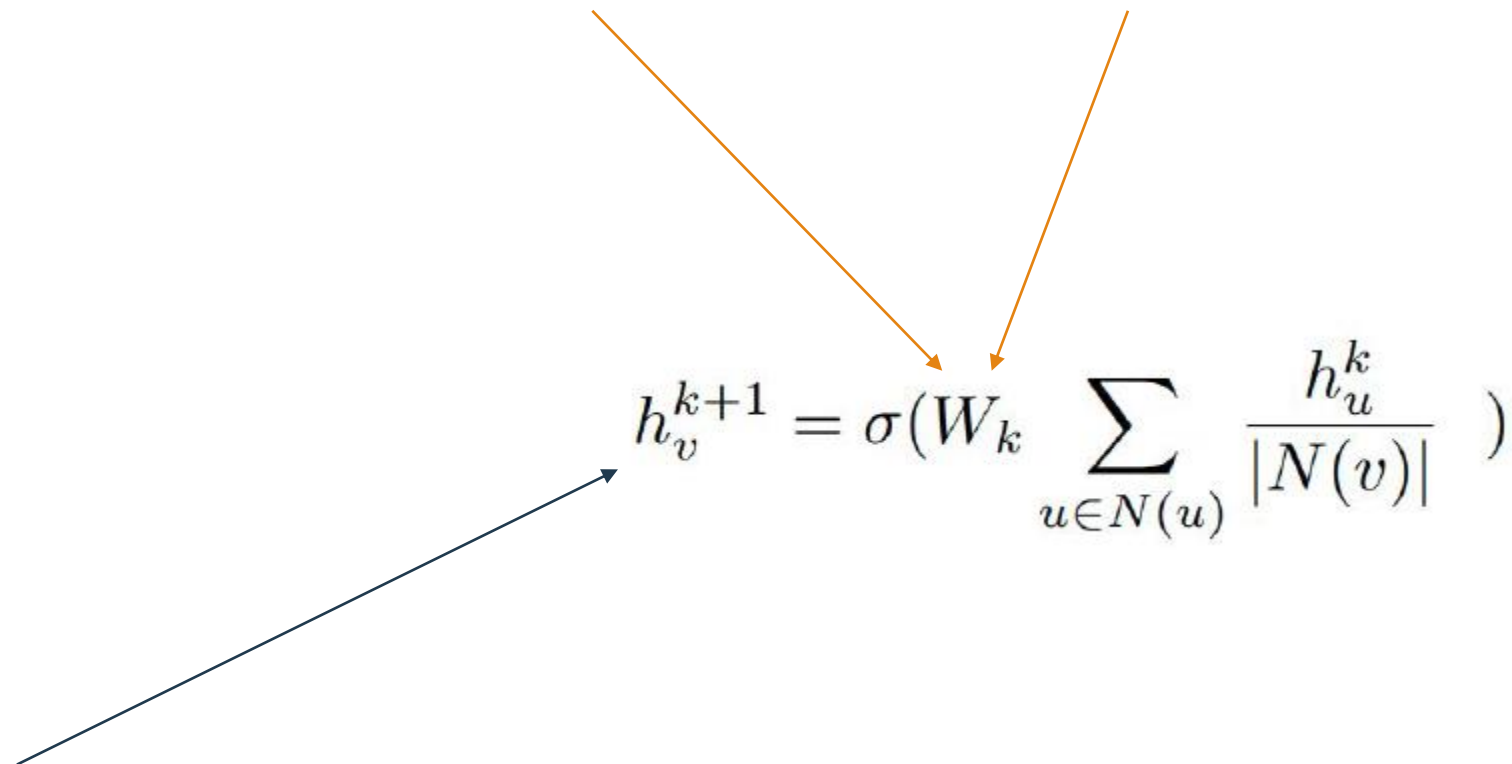
conv1 = GCNConv(num_of_input_features, num_of_out_features) # init of the GNN
```


$$h_v^{k+1} = \sigma \left( W_k \sum_{u \in N(v)} \frac{h_u^k}{|N(v)|} \right)$$

## Pytorch Geometric

```
import torch
from torch.nn import Linear
from torch_geometric.nn import GCNConv
```

```
conv1 = GCNConv(num_of_input_features, num_of_out_features) # init of the GNN
```


$$h_v^{k+1} = \sigma \left( W_k \sum_{u \in N(v)} \frac{h_u^k}{|N(v)|} \right)$$

```
h = conv1(input_features, graph_edges) # running the GNN on an input
```

## Pytorch Geometric

```
import torch
from torch.nn import Linear
from torch_geometric.nn import GCNConv
```

```
conv1 = GCNConv(num_of_input_features, num_of_out_features) # init of the GNN
```

$$h_v^{k+1} = \sigma \left( W_k \sum_{u \in N(v)} \frac{h_u^k}{|N(v)|} \right)$$

```
h = conv1(input_features, graph_edges) # running the GNN on an input
```

## Pytorch Geometric

```
import torch
from torch.nn import Linear
from torch_geometric.nn import GCNConv

class GCN(torch.nn.Module):
    def __init__(self):
        super().__init__()
        torch.manual_seed(1234)
        self.conv1 = GCNConv(dataset.num_features, 4)
        self.conv2 = GCNConv(4, 4)
        self.conv3 = GCNConv(4, 2)
        self.classifier = Linear(2, dataset.num_classes)

    def forward(self, x, edge_index):
        h = self.conv1(x, edge_index)
        h = h.tanh()
        h = self.conv2(h, edge_index)
        h = h.tanh()
        h = self.conv3(h, edge_index)
        h = h.tanh() # Final GNN embedding space.

        # Apply a final (linear) classifier.
        out = self.classifier(h)

    return out, h
```

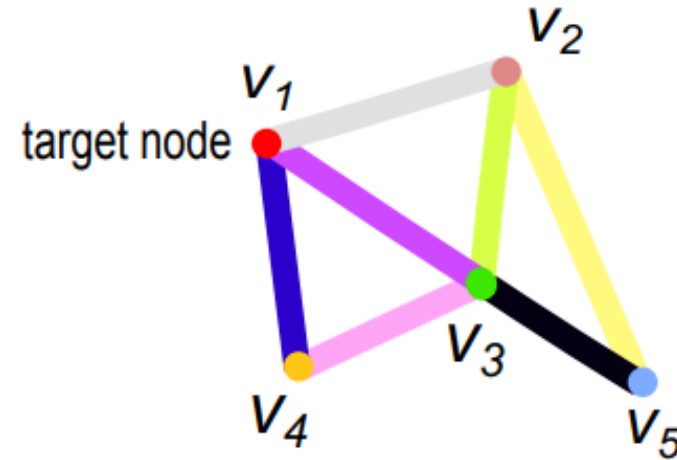
## Pytorch Geometric

```
model = GCN()
criterion = torch.nn.CrossEntropyLoss() # Define loss criterion.
optimizer = torch.optim.Adam(model.parameters(), lr=0.01) # Define optimizer.

def train(data):
    optimizer.zero_grad() # Clear gradients.
    out, h = model(data.x, data.edge_index) # Perform a single forward pass.
    loss = criterion(out[data.train_mask], data.y[data.train_mask]) # Compute the loss solely based on the training nodes.
    loss.backward() # Derive gradients.
    optimizer.step() # Update parameters based on gradients.
    return loss, h
```

# Message Passing Neural Networks

1. A graph  $G = (V, E)$ .
2. For each node  $i \in V$  we have an initial vector  $h_i^{(0)} \in \mathbb{R}^{l_0}$ .



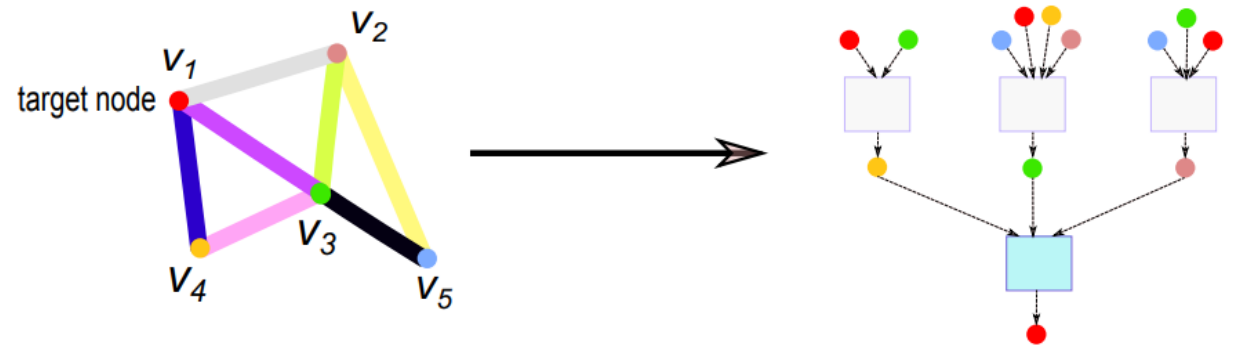


# Message Passing Neural Networks

1. A graph  $G = (V, E)$ .
2. For each node  $i \in V$  we have an initial vector  $h_i^{(0)} \in \mathbb{R}^{l_0}$ .

Given the above data, the feedforward neural algorithm on  $G$  executes  $L$  message passing schemes defined recursively for  $0 \leq k \leq L$  by:

$$h_i^{(k)} := \alpha^k \left( h_i^{(k-1)}, E_{j \in \mathcal{N}(i)} \left( \phi^k(h_i^{(k-1)}, h_j^{(k-1)}, e_{i,j}) \right) \right) \in \mathbb{R}^{l_k},$$



# Message Passing Neural Networks

1. A graph  $G = (V, E)$ .
2. For each node  $i \in V$  we have an initial vector  $h_i^{(0)} \in \mathbb{R}^{l_0}$ .

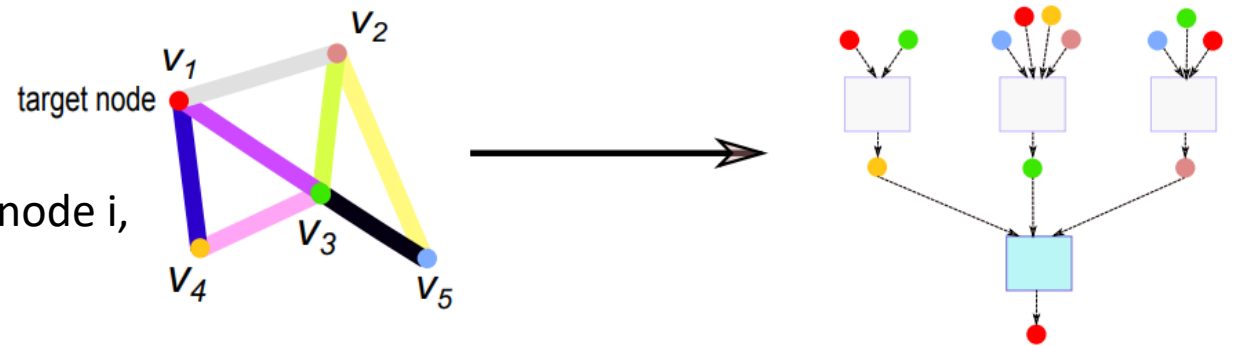
Given the above data, the feedforward neural algorithm on  $G$  executes  $L$  message passing schemes defined recursively for  $0 \leq k \leq L$  by:

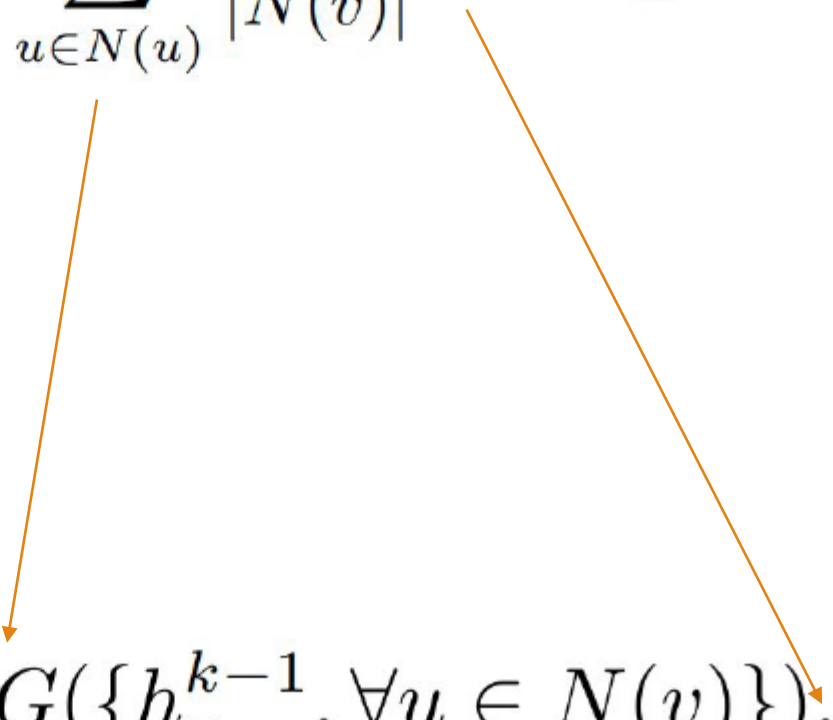
$$h_i^{(k)} := \alpha^k \left( h_i^{(k-1)}, E_{j \in \mathcal{N}(i)} \left( \phi^k (h_i^{(k-1)}, h_j^{(k-1)}, e_{i,j}) \right) \right) \in \mathbb{R}^{l_k},$$

where  $e_{ij} \in \mathbb{R}^D$  is an edge feature from the node  $j$  to the node  $i$ ,

$E$  is a permutation invariant differentiable function,

$\alpha^k, \phi^k$  are trainable differentiable functions (regular neural networks)



$$h_v^{k+1} = \sigma\left(W_k \sum_{u \in N(v)} \frac{h_u^k}{|N(v)|} + B_k h_v^k\right)$$


The diagram consists of two orange arrows. The first arrow originates from the summation term  $\sum_{u \in N(v)} \frac{h_u^k}{|N(v)|}$  in the first equation and points to the  $AGG(\{h_u^{k-1}, \forall u \in N(v)\})$  term in the second equation. The second arrow originates from the bias term  $B_k h_v^k$  in the first equation and points to the  $B_k h_v^k$  term in the second equation.

$$h_v^{k+1} = \sigma([W_k \cdot AGG(\{h_u^{k-1}, \forall u \in N(v)\}), B_k h_v^k])$$

$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \alpha_{vu} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)} \right)$$

Attention weights

**Not all node's neighbors are equally important**

**Attention** is inspired by cognitive attention.

The **attention**  $\alpha_{vu}$  focuses on the important parts of the input data and fades out the rest.

**Idea:** the NN should devote more computing power on that small but important part of the data. Which part of the data is more important depends on the context and is learned through training.

$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \boxed{\alpha_{vu}} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)} \right)$$

Attention weights

**Goal:** Specify **arbitrary importance** to different neighbors of each node in the graph

**Idea:** Compute embedding  $\mathbf{h}_v^{(l)}$  of each node in the graph following an **attention**

**strategy:** ✦ Nodes attend over their neighborhoods' message

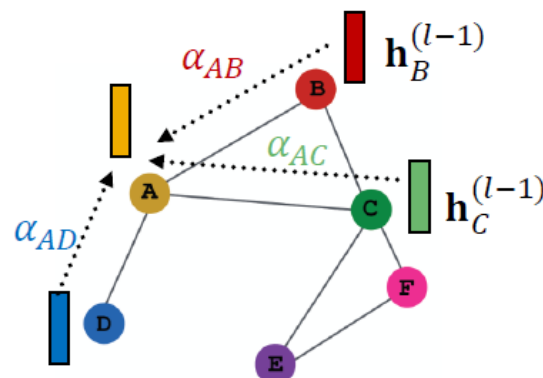
✦ Implicitly specifying different weights to different nodes in a neighborhood

$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \boxed{\alpha_{vu}} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)} \right)$$

Attention weights

- Let  $\alpha_{vu}$  be computed as a byproduct of an **attention mechanism**  $a$ :
  - (1) Let  $a$  compute **attention coefficients**  $e_{vu}$  across pairs of nodes  $u, v$  based on their messages:

$$e_{vu} = a(\mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}, \mathbf{W}^{(l)} \mathbf{h}_v^{(l-1)})$$



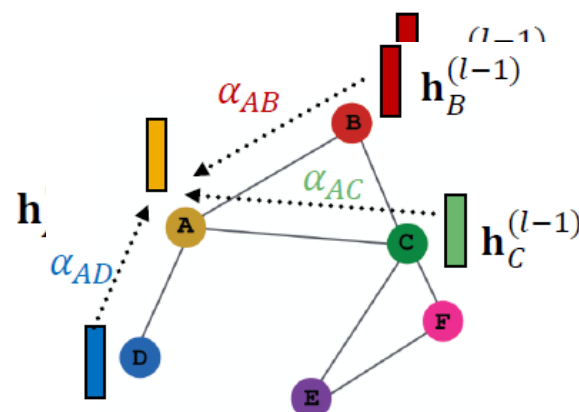
$$\mathbf{h}_v^{(l)} = \sigma\left(\sum_{u \in N(v)} \boxed{\alpha_{vu}} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}\right)$$

Attention weights

- Let  $\alpha_{vu}$  be computed as a byproduct of an **attention mechanism**  $a$ :
  - (1) Let  $a$  compute **attention coefficients**  $e_{vu}$  across pairs of nodes  $u, v$  based on their messages:

$$e_{vu} = a(\mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}, \mathbf{W}^{(l)} \mathbf{h}_v^{(l-1)})$$

- $e_{vu}$  indicates the importance of  $u$ 's message to node  $v$



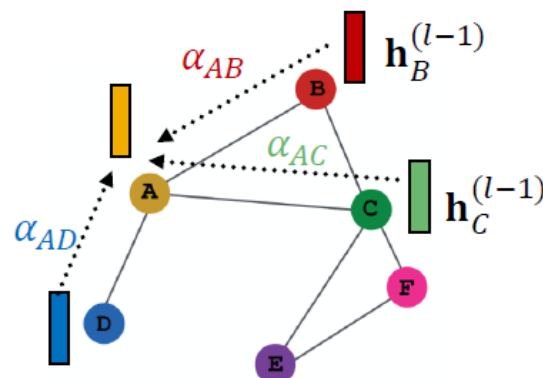
$$e_{AB} = a(\mathbf{W}^{(l)} \mathbf{h}_A^{(l-1)}, \mathbf{W}^{(l)} \mathbf{h}_B^{(l-1)})$$

$$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \boxed{\alpha_{vu}} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)} \right)$$

Attention weights

- **Normalize**  $e_{vu}$  into the **final attention weight**  $\alpha_{vu}$ 
  - Use the **softmax** function, so that  $\sum_{u \in N(v)} \alpha_{vu} = 1$ :

$$\alpha_{vu} = \frac{\exp(e_{vu})}{\sum_{k \in N(v)} \exp(e_{vk})}$$





$$\mathbf{h}_v^{(l)} = \sigma\left(\sum_{u \in N(v)} \boxed{\alpha_{vu}} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}\right)$$

Attention weights

- **Normalize**  $e_{vu}$  into the **final attention weight**  $\alpha_{vu}$ 
  - Use the **softmax** function, so that  $\sum_{u \in N(v)} \alpha_{vu} = 1$ :

$$\alpha_{vu} = \frac{\exp(e_{vu})}{\sum_{k \in N(v)} \exp(e_{vk})}$$

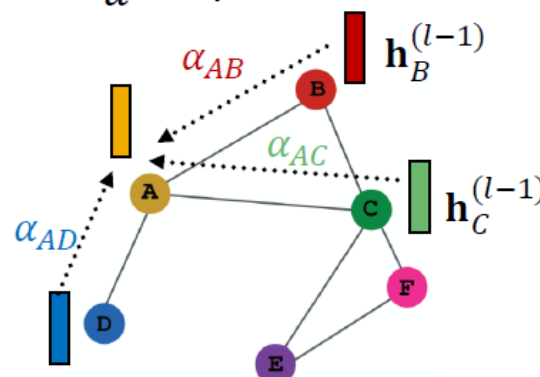
- **Weighted sum** based on the **final attention weight**

$\alpha_{vu}$

$$\mathbf{h}_v^{(l)} = \sigma\left(\sum_{u \in N(v)} \alpha_{vu} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}\right)$$

**Weighted sum using  $\alpha_{AB}$ ,  $\alpha_{AC}$ ,  $\alpha_{AD}$ :**

$$\mathbf{h}_A^{(l)} = \sigma(\alpha_{AB} \mathbf{W}^{(l)} \mathbf{h}_B^{(l-1)} + \alpha_{AC} \mathbf{W}^{(l)} \mathbf{h}_C^{(l-1)} + \alpha_{AD} \mathbf{W}^{(l)} \mathbf{h}_D^{(l-1)})$$



## The over smoothing problem

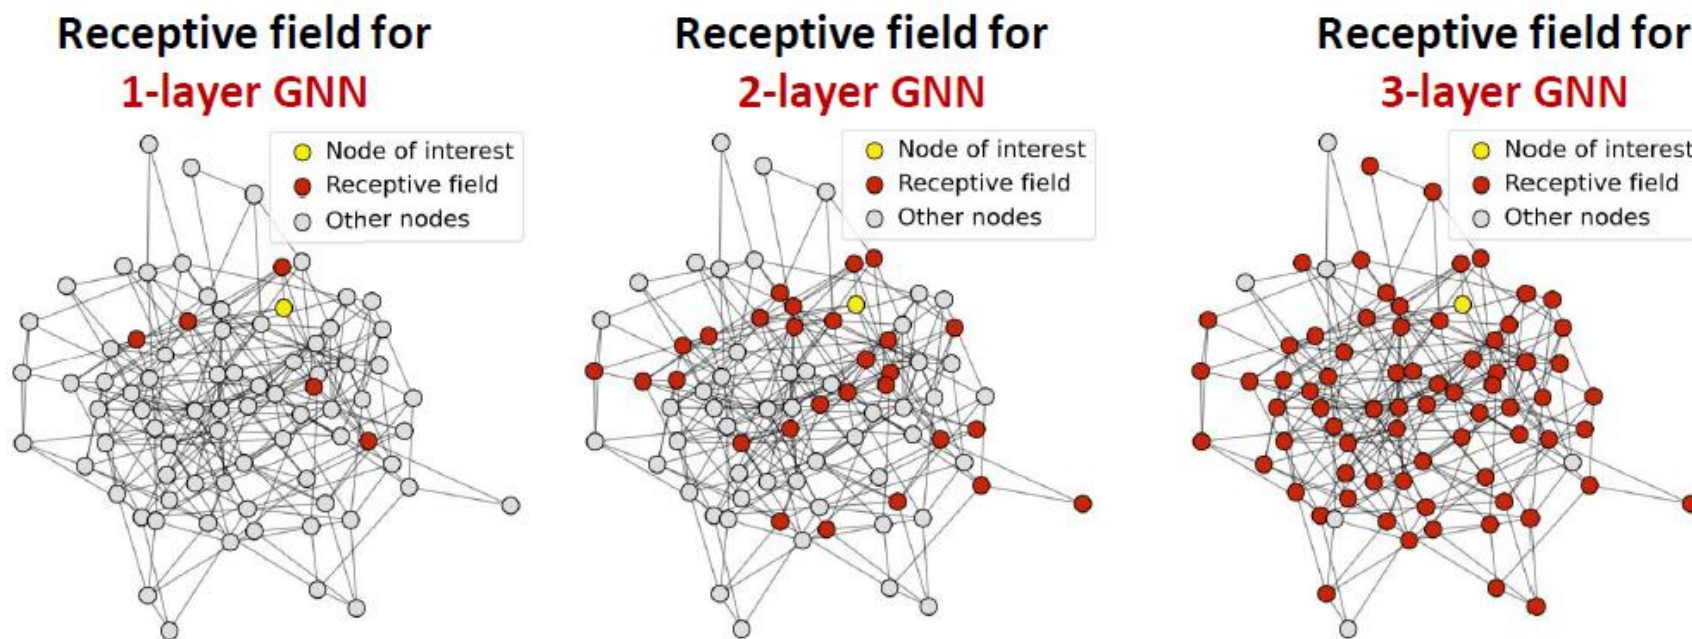
### The Issue of stacking many GNN layers

- ✦ GNN suffers from **the over-smoothing problem**
- **The over-smoothing problem:** all the node embeddings converge to the same value
- ✦ This is bad because we **want to use node embeddings to differentiate nodes**
- **Why does the over-smoothing problem happen?**

## The over smoothing problem

**Receptive field:** the set of nodes that determine the embedding of a node of interest

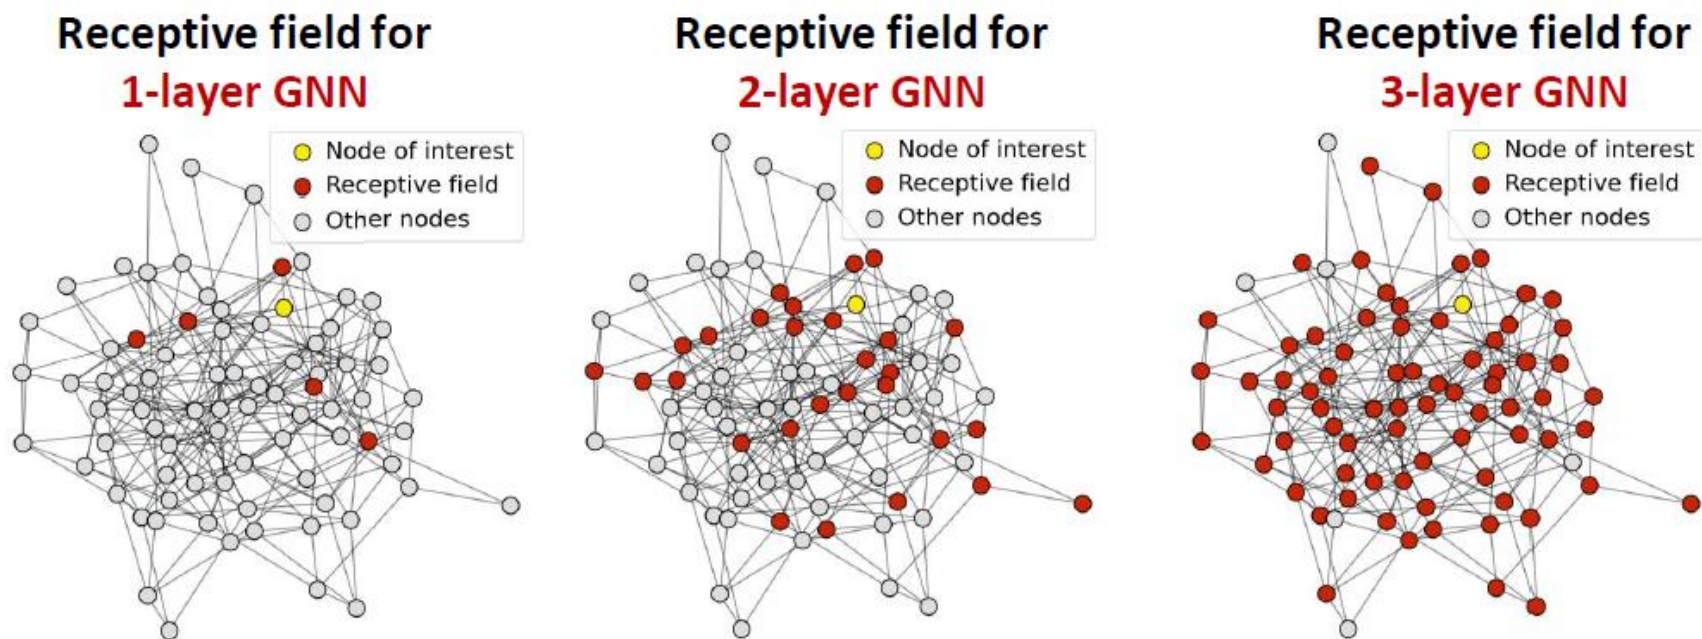
★ In a  $K$ -layer GNN, each node has a receptive field of  $K$ -hop neighborhood



# The over smoothing problem

Receptive field overlap for two nodes

The shared neighbors quickly grows when we increase the number of hops (num of GNN layers)



## The over smoothing problem

We can explain over-smoothing via the notion of receptive field

- ★ We knew the embedding of a node is determined by its receptive field
- ★ If two nodes have highly-overlapped receptive fields, then their embeddings are highly similar
- ★ Stack many GNN layers -> nodes will have highly-overlapped receptive fields -> node embeddings will be highly similar -> suffer from the over-smoothing problem
- Next: how do we overcome over-smoothing problem?

## The over smoothing problem

### What do we learn from the over-smoothing problem? ■

#### Lesson 1: Be cautious when adding GNN layers

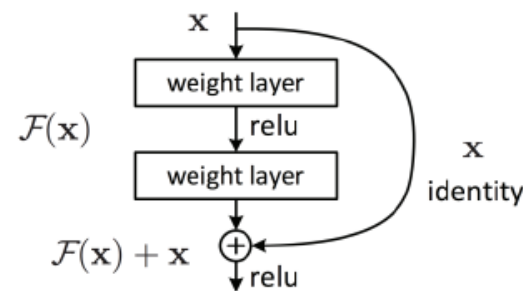
- ✦ Unlike neural networks in other domains (CNN for image classification), **adding more GNN layers do not always help**
- ✦ **Step 1: Analyze the necessary receptive field** to solve your problem. E.g., by computing the diameter of the graph
- ✦ **Step 2:** Set number of GNN layers  $L$  to be a bit more than the receptive field we like. **Do not set  $L$  to be unnecessarily large!**

## One possible solution

- A standard GCN layer

- $$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \mathbf{W}^{(l)} \frac{\mathbf{h}_u^{(l-1)}}{|N(v)|} \right)$$

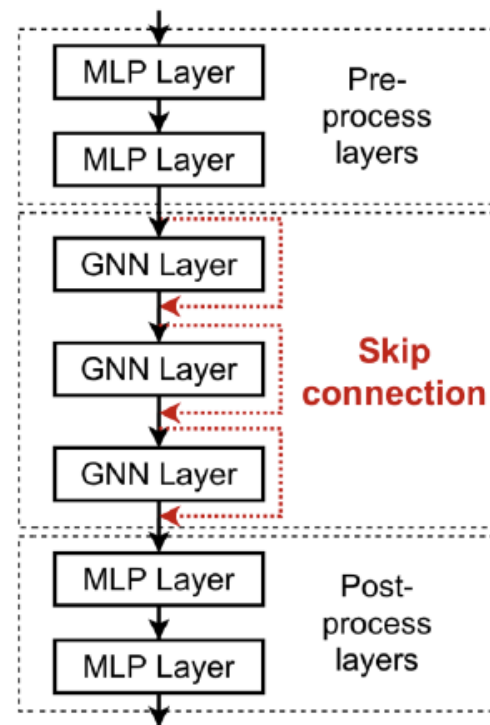
This is our  $F(\mathbf{x})$



- A GCN layer with skip connection

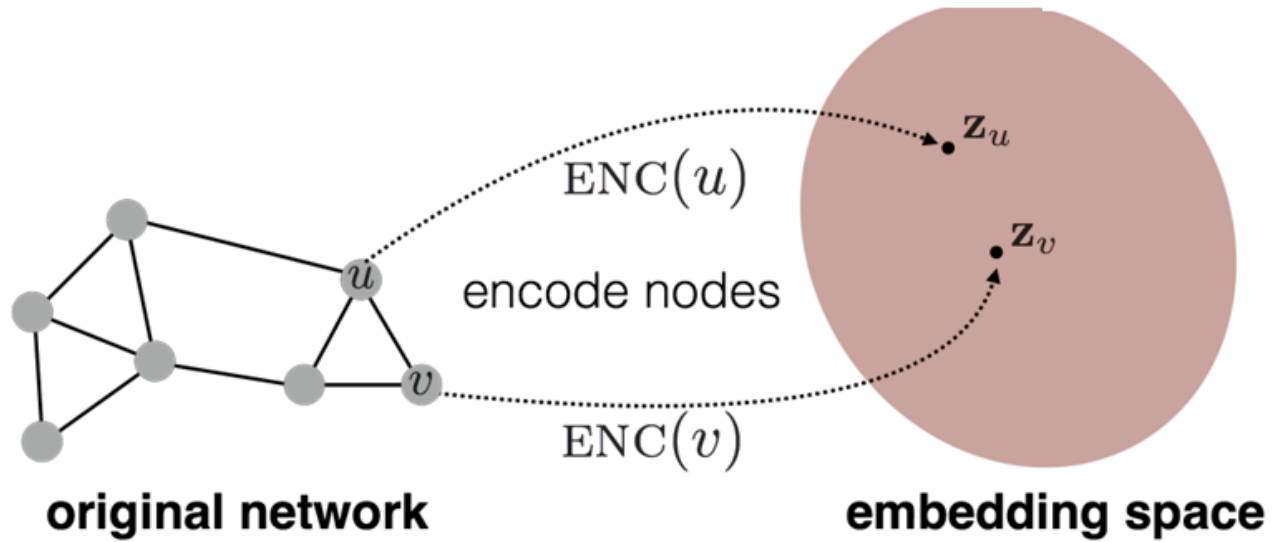
- $$\mathbf{h}_v^{(l)} = \sigma \left( \sum_{u \in N(v)} \mathbf{W}^{(l)} \frac{\mathbf{h}_u^{(l-1)}}{|N(v)|} + \mathbf{h}_v^{(l-1)} \right)$$

$F(\mathbf{x})$                       +                       $\mathbf{x}$



## Graph Representation learning

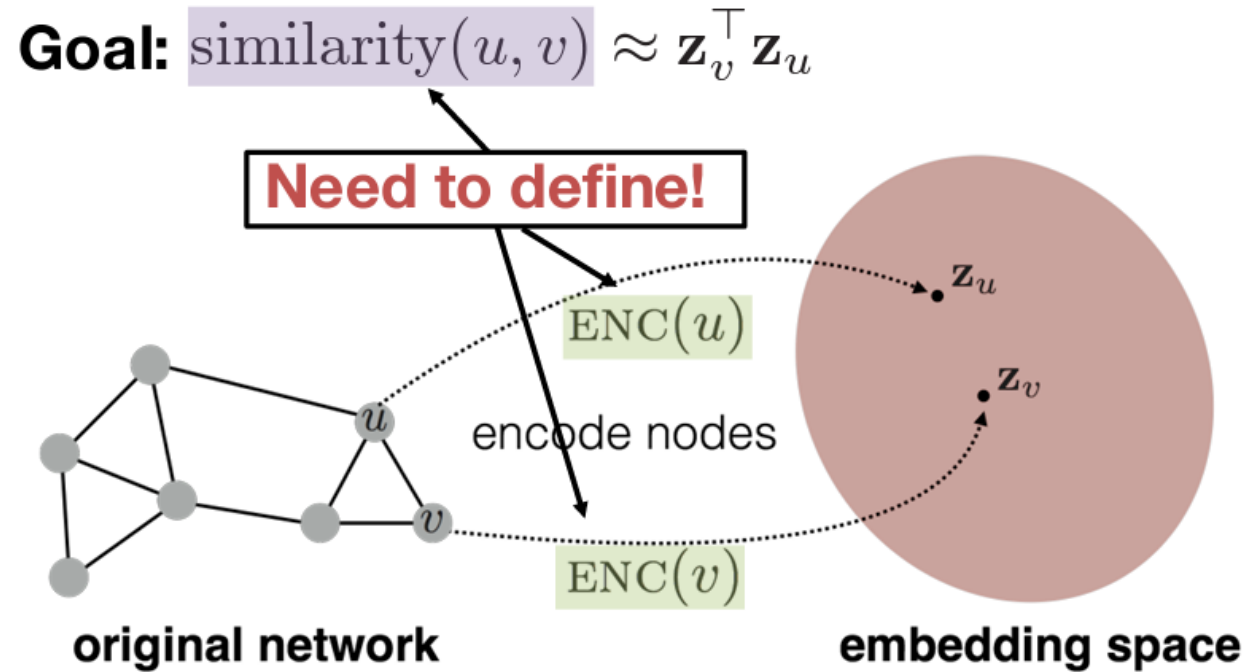
**Goal** is to encode nodes so that similarity in the embedding space (e.g., dot product) approximates similarity in the original network.





## Graph Representation learning

**Goal** is to encode nodes so that similarity in the embedding space (e.g., dot product) approximates similarity in the original network.



## Graph Representation learning

- **Encoder** maps each node to a low-dimensional vector.

$$\text{ENC}(v) = \mathbf{z}_v$$

node in the input graph

d-dimensional embedding

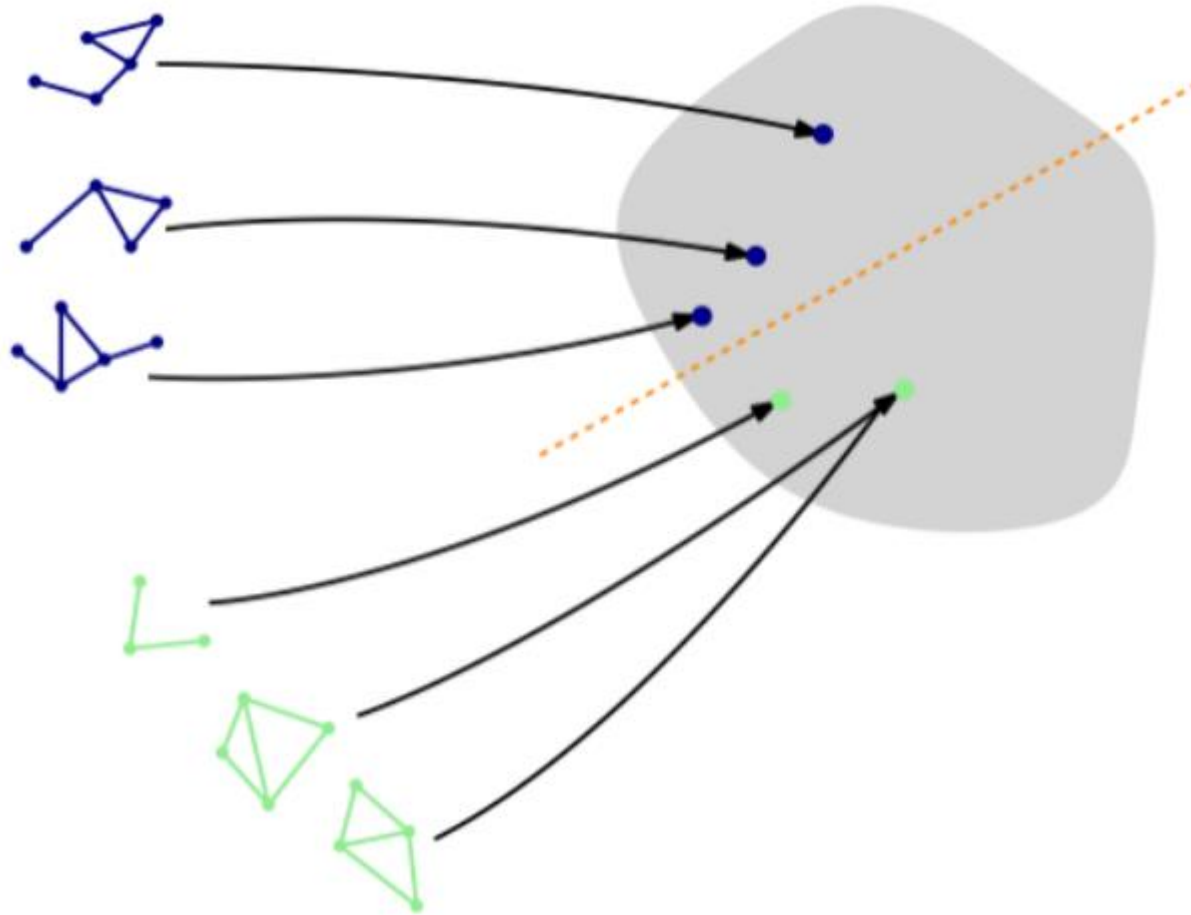
- **Similarity function** specifies how relationships in vector space map to relationships in the original network.

$$\text{similarity}(u, v) \approx \mathbf{z}_v^T \mathbf{z}_u$$

Similarity of  $u$  and  $v$  in the original network

dot product between node embeddings

## Graph Classification with Pytorch Geometric



[https://colab.research.google.com/drive/1I8a0DfQ3fl7Njc62\\_mVXUlcAleUclnb?usp=sharing](https://colab.research.google.com/drive/1I8a0DfQ3fl7Njc62_mVXUlcAleUclnb?usp=sharing)

## Graph Classification with Pytorch Geometric

Typically we have the following steps :

- (1) Embed each node of the graph using one of the GNNs we introduced here.
- (2) Aggregate all node embeddings and define a global graph embedding using aggregation invariant function (sum,min,max). For instance :

$$\mathbf{x}_{\mathcal{G}} = \frac{1}{|\mathcal{V}|} \sum_{v \in \mathcal{V}} x_v^{(L)}$$

[https://colab.research.google.com/drive/1I8a0DfQ3fl7Njc62\\_mVXUlcAleUclnb?usp=sharing](https://colab.research.google.com/drive/1I8a0DfQ3fl7Njc62_mVXUlcAleUclnb?usp=sharing)