

# Introduction to neural networks II

# Objectives

- Recall the feedforward of a neural network
- The binary classification problem
- How to build a neural network that can classify a binary labeled data?
- How can we build a neural network that can classify a multi-labeled data?
- Introduction of the sigmoid function
- Introduction of the softmax function
- Computational Graphs
- Backprop and automatic differentiation
- The approximation power of neural networks (universal approximation theorem)

# How do we use a neural network as a classifier?

Last time we learned the following:

- The building block of a neural network

# How do we use a neural network as a classifier?

Last time we learned the following:

- The building block of a neural network
- How to build a neural network.

# How do we use a neural network as a classifier?

Last time we learned the following:

- The building block of a neural network
- How to build a neural network.
- Neural network is essentially a mathematical function  $f: R^n \rightarrow R^m$ .

# How do we use a neural network as a classifier?

Last time we learned the following:

- The building block of a neural network
- How to build a neural network.
- Neural network is essentially a mathematical function  $f: R^n \rightarrow R^m$ .
- Given an input  $x$ , how to feedforward  $x$  through a neural network and obtain an output  $f(x)$

# How do we use a neural network as a classifier?

Last time we learned the following:

- The building block of a neural network
- How to build a neural network.
- Neural network is essentially a mathematical function  $f: R^n \rightarrow R^m$ .
- Given an input  $x$ , how to feedforward  $x$  through a neural network and obtain an output  $f(x)$
- How to train a neural network :
  - Define a cost function
  - For each example in the training set feedforward that example and compute the error
  - Use backpropagation to adjust the weights of the network so that it behaves better with respect to the input example

# How do we use a neural network as a classifier?

Last time we learned the following:

- The building block of a neural network
- How to build a neural network.
- Neural network is essentially a mathematical function  $f: R^n \rightarrow R^m$ .
- Given an input  $x$ , how to feedforward  $x$  through a neural network and obtain an output  $f(x)$
- How to train a neural network :
  - Define a cost function
  - For each example in the training set feedforward that example and compute the error
  - Use backpropagation to adjust the weights of the network so that it behaves better with respect to the input example

Lets recall the feedforward algorithm before first.



# Feedforward Neural Network

How do we compute a feedforward neural network on an input  $x$  ?

# Feedforward Neural Network

Start with an input  $x = a^{(0)}$ . In the picture, this is represented by the first layer of nodes. We will call this layer 0.

$$x = a^{(0)}$$

# Feedforward Neural Network

We apply the weight  $W^{(1)}$  coming from the edges between layer 0 and layer 1 and add the biases and then apply the Activation function on the resulting vector coordinate-wise.

$$x = a^{(0)} \longrightarrow \sigma(W^{(1)}a^{(0)} + b^{(1)})$$

$W^{(1)}$  : Edges between  
layer 0 and layer 1

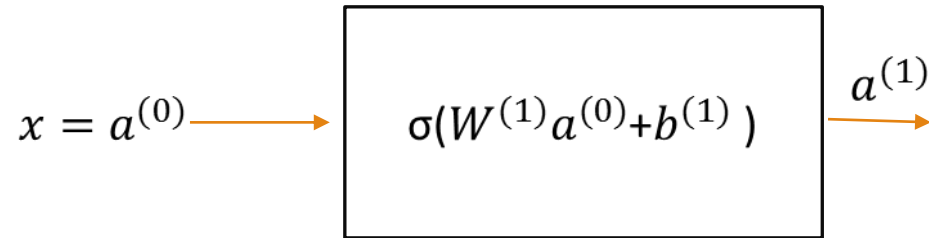
$a^{(0)}$  : input

$b^{(1)}$  : biases applied to layer 1

$\sigma$  : activation function

# Feedforward Neural Network

We will call the output of this computation  $a^{(1)}$ . This is now represented by the nodes in layer 1.



$W^{(1)}$  : Edges between  
layer 0 and layer 1

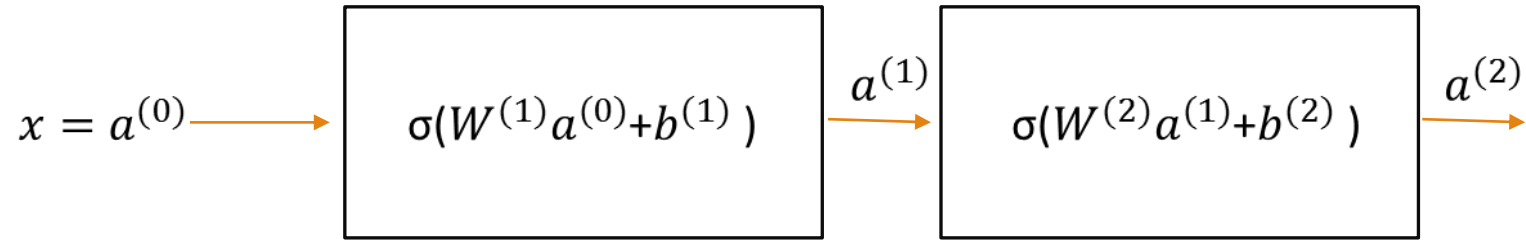
$a^{(0)}$  : input

$b^{(1)}$  : biases applied to layer 1

$\sigma$  : activation function

# Feedforward Neural Network

Repeat.



$W^{(2)}$  : Edges between  
layer 1 and layer 2

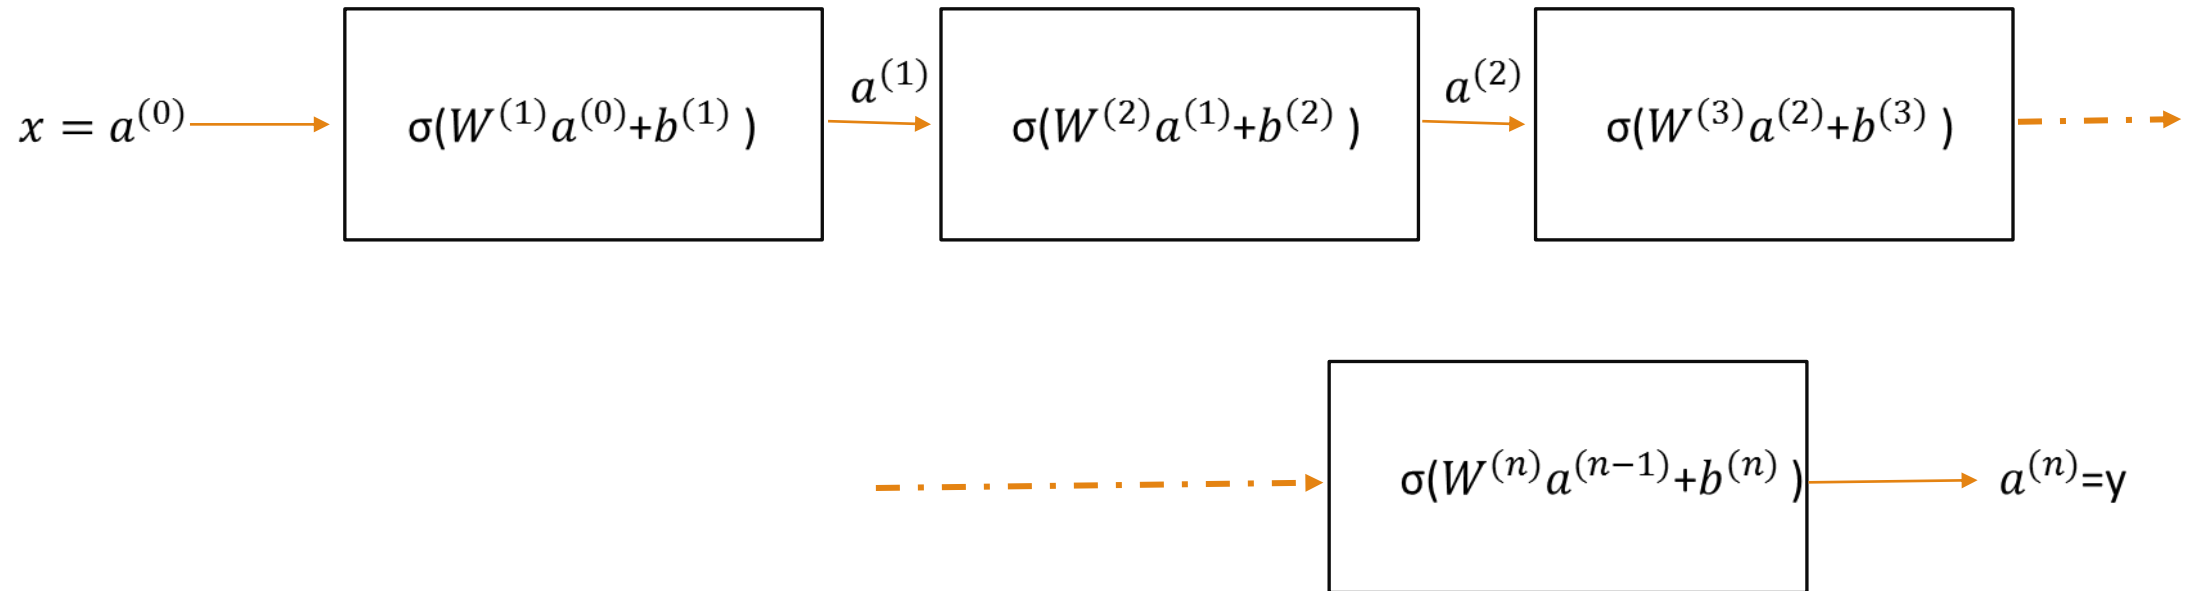
$a^{(1)}$  : input from layer 1

$b^{(2)}$  : biases applied to layer 2

$\sigma$  : activation function

# Feedforward Neural Network

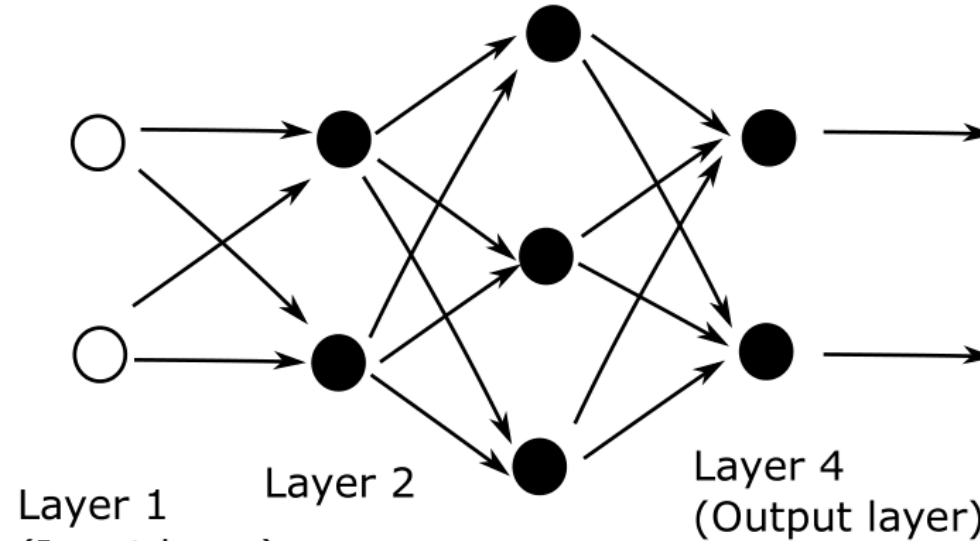
Until you finish the neural network and get the final output.



# Example

We will use an example from [this](#) paper.


(note that the convention of the index is a little different here)



# Example

We will use an example from [this](#) paper.

(note that the convention of the index is a little different here)

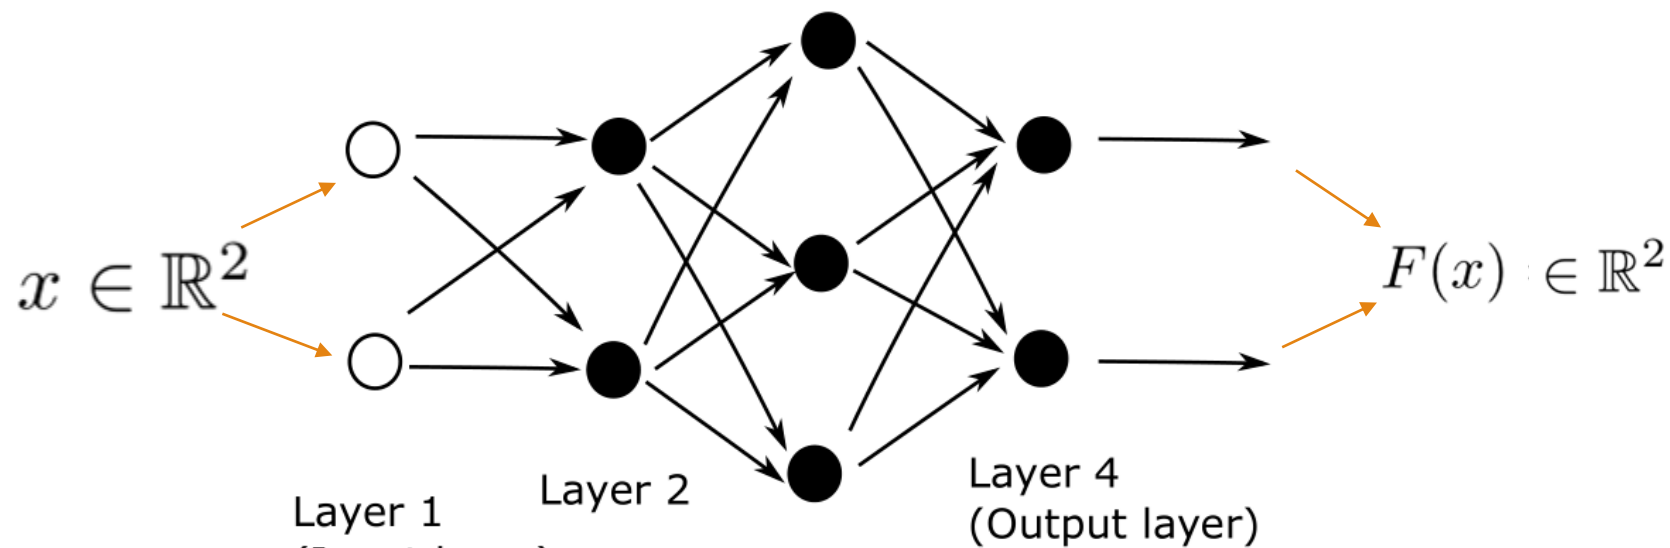




# Example

We will use an example from [this](#) paper.

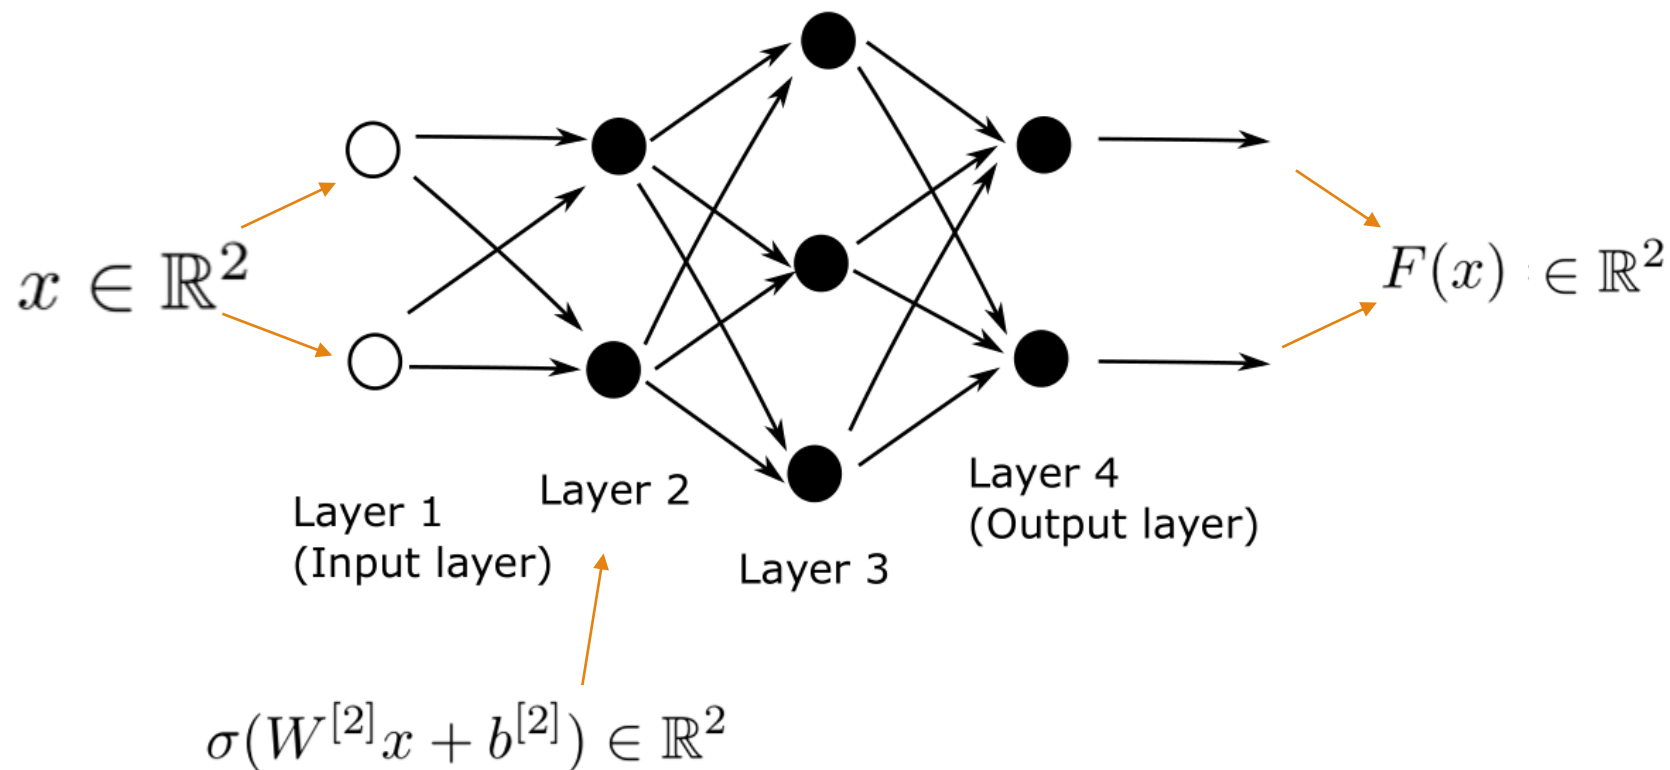
(note that the convention of the index is a little different here)



# Example

We will use an example from [this](#) paper.

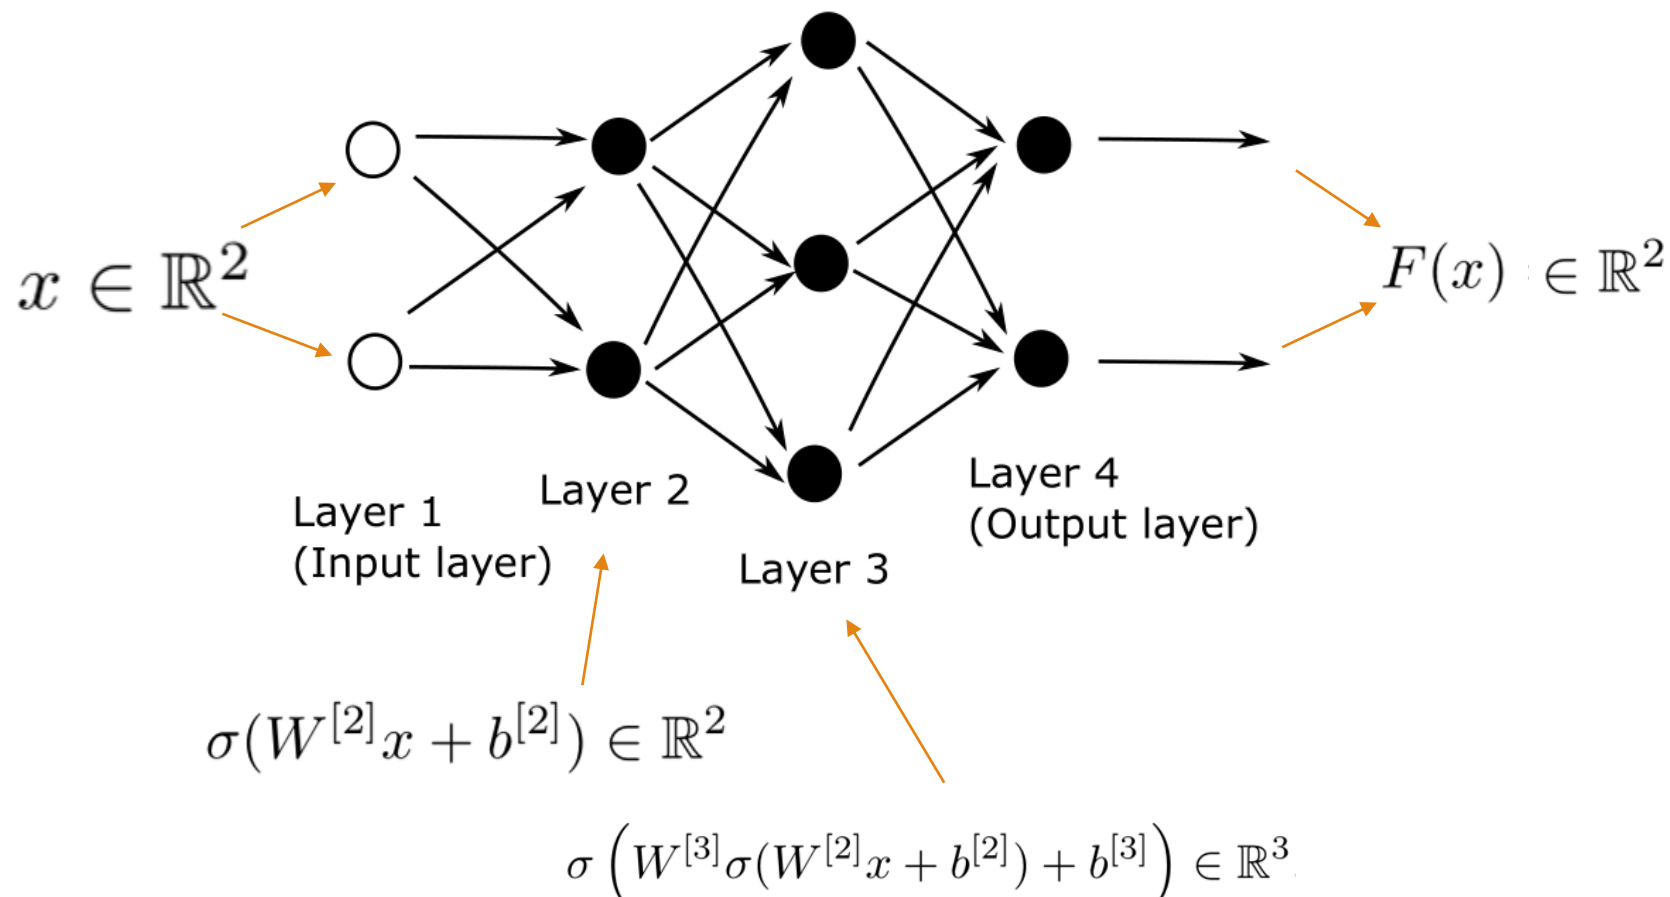
(note that the convention of the index is a little different here)



# Example

We will use an example from [this](#) paper.

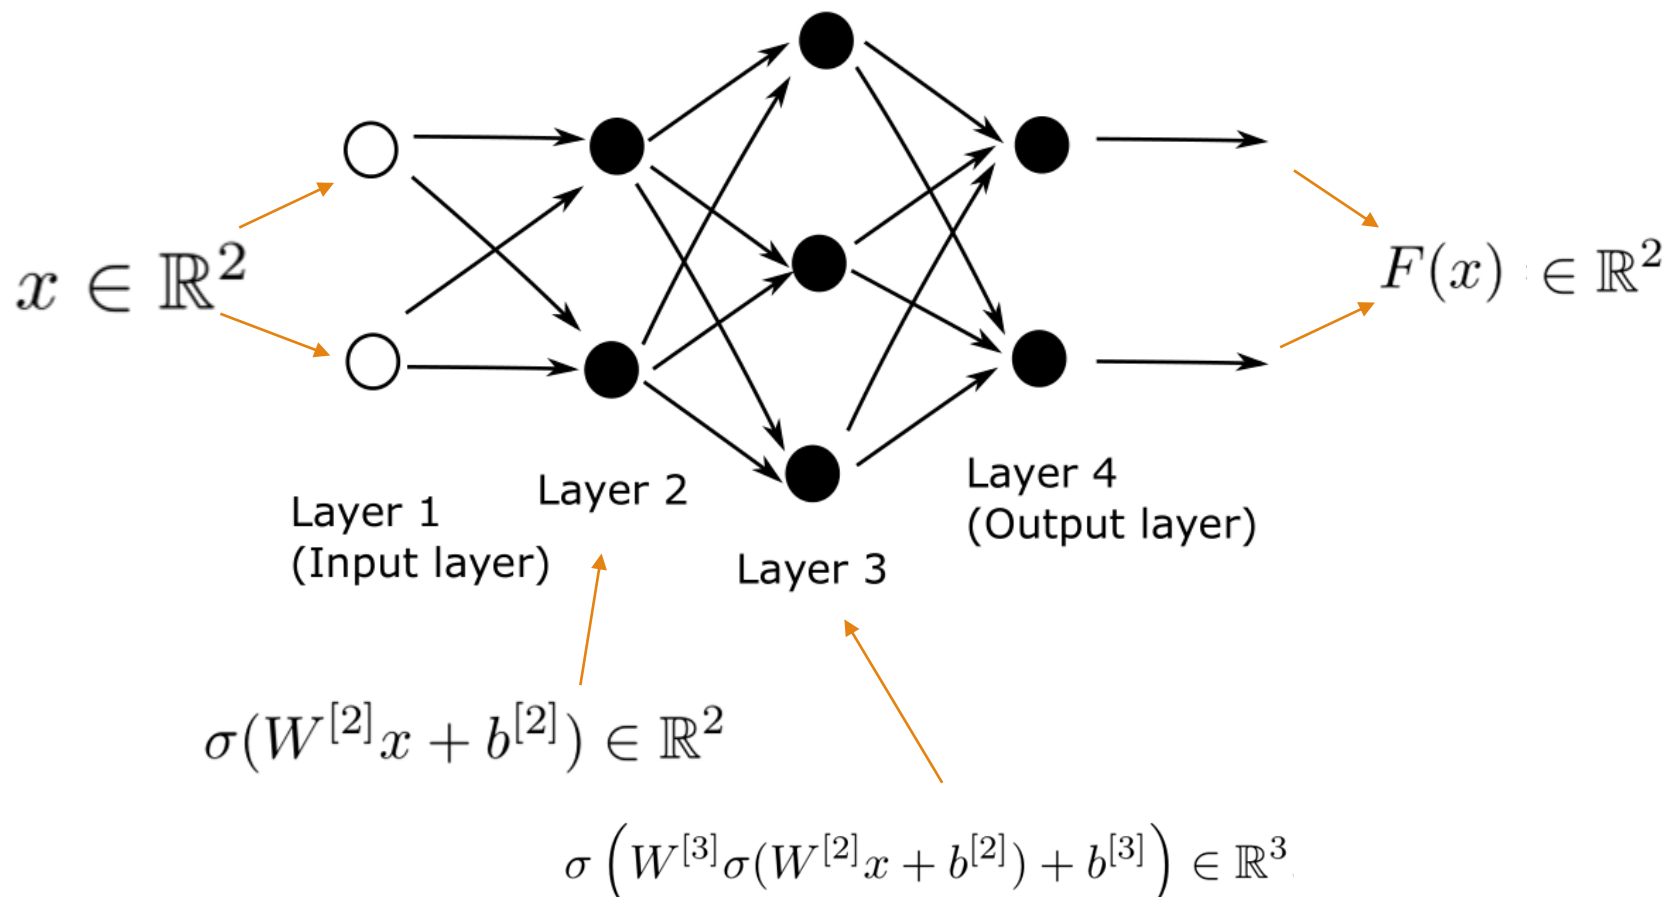
(note that the convention of the index is a little different here)



# Example

We will use an example from [this](#) paper.

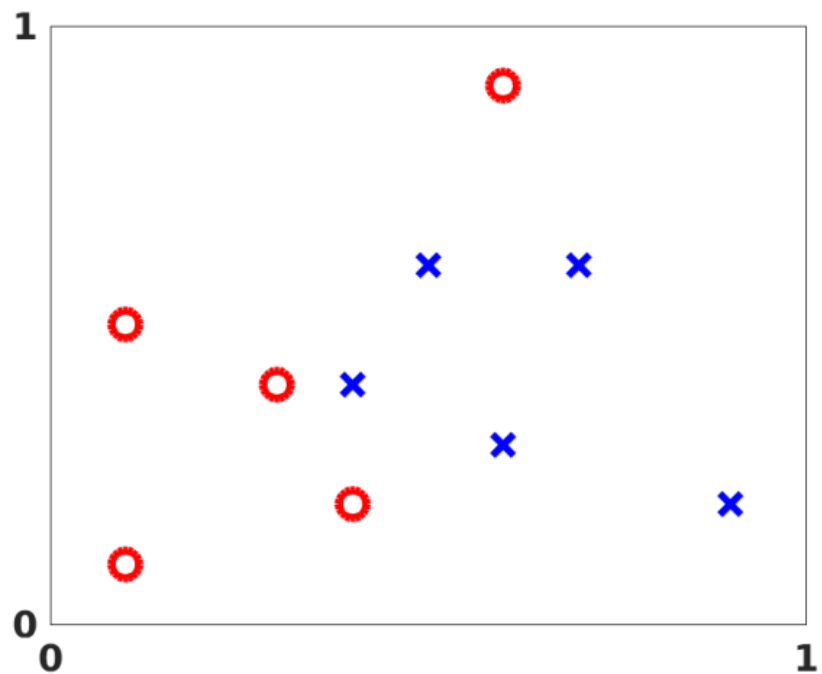
(note that the convention of the index is a little different here)



Final function representing the neural network

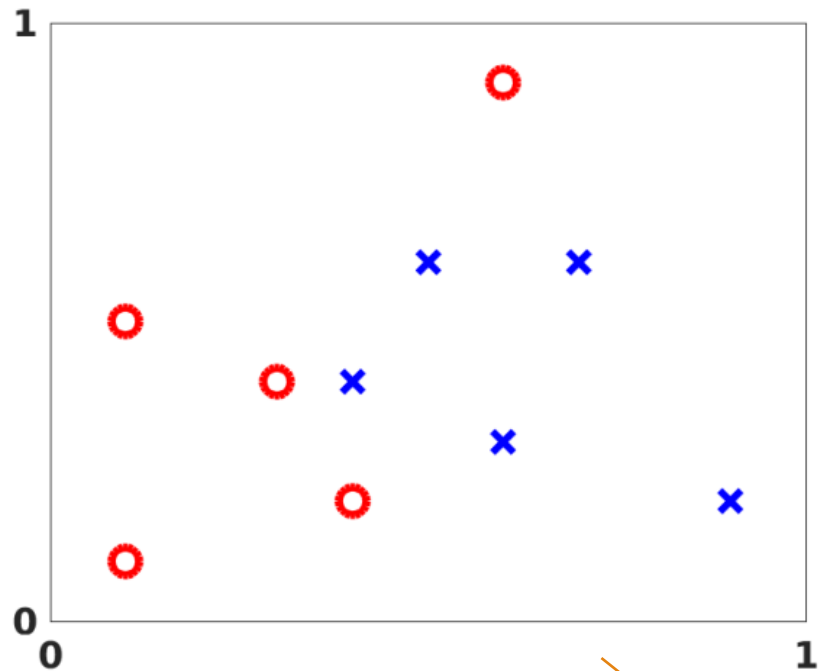
$$F(x) = \sigma \left( W^{[4]} \sigma \left( W^{[3]} \sigma(W^{[2]}x + b^{[2]}) + b^{[3]} \right) + b^{[4]} \right) \in \mathbb{R}^2.$$

# Example



Input : labeled data X

# Example

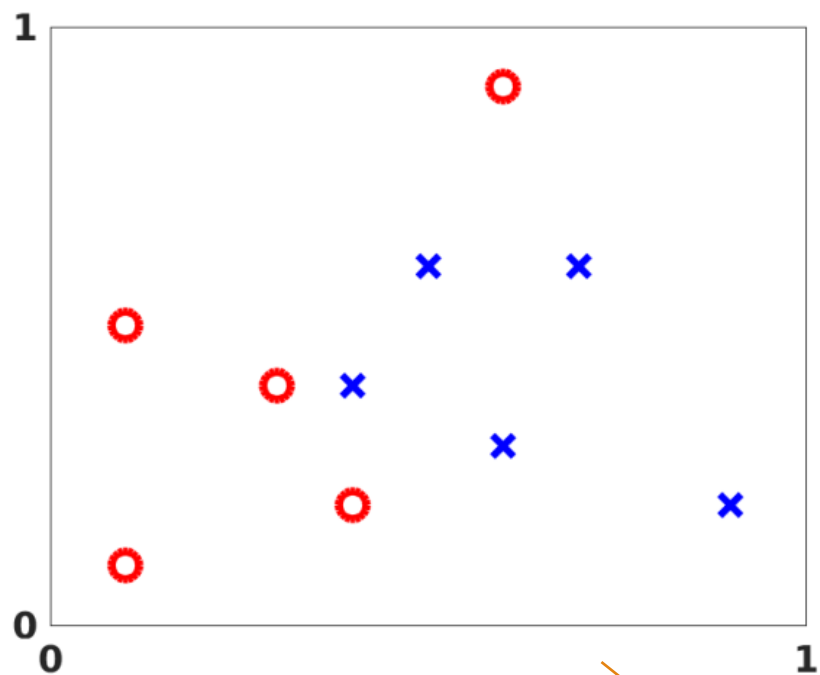


Input : labeled data X


$$\text{Cost} \left( W^{[2]}, W^{[3]}, W^{[4]}, b^{[2]}, b^{[3]}, b^{[4]} \right) = \frac{1}{10} \sum_{i=1}^{10} \frac{1}{2} \|y(x^{\{i\}}) - F(x^{\{i\}})\|_2^2.$$

the difference between the output given by the network and the actual label

# Example



Input : labeled data X



Minimize the cost function

$$\text{Cost} \left( W^{[2]}, W^{[3]}, W^{[4]}, b^{[2]}, b^{[3]}, b^{[4]} \right) = \frac{1}{10} \sum_{i=1}^{10} \frac{1}{2} \|y(x^{\{i\}}) - F(x^{\{i\}})\|_2^2.$$

the difference between the output given by the network and the actual label

# Binary classification

Now suppose that we have data set that consists of images of cats and dogs and we built a neural network that takes as input an image from this data set and gives out a vector in  $\mathbb{R}^1$  (a real number).

How exactly do we use this vector for our classification task ? In general the output  $f(x)$  coming from the neural network Does not match the class  $\{\pm 1\}$  of the input point  $x$  (it could be any real number).





# Binary classification

This function takes a tensor of size `input_size` and returns a real number.

How can we constrain the output to be between -1 and +1?

```
import torch
import torch.nn as nn

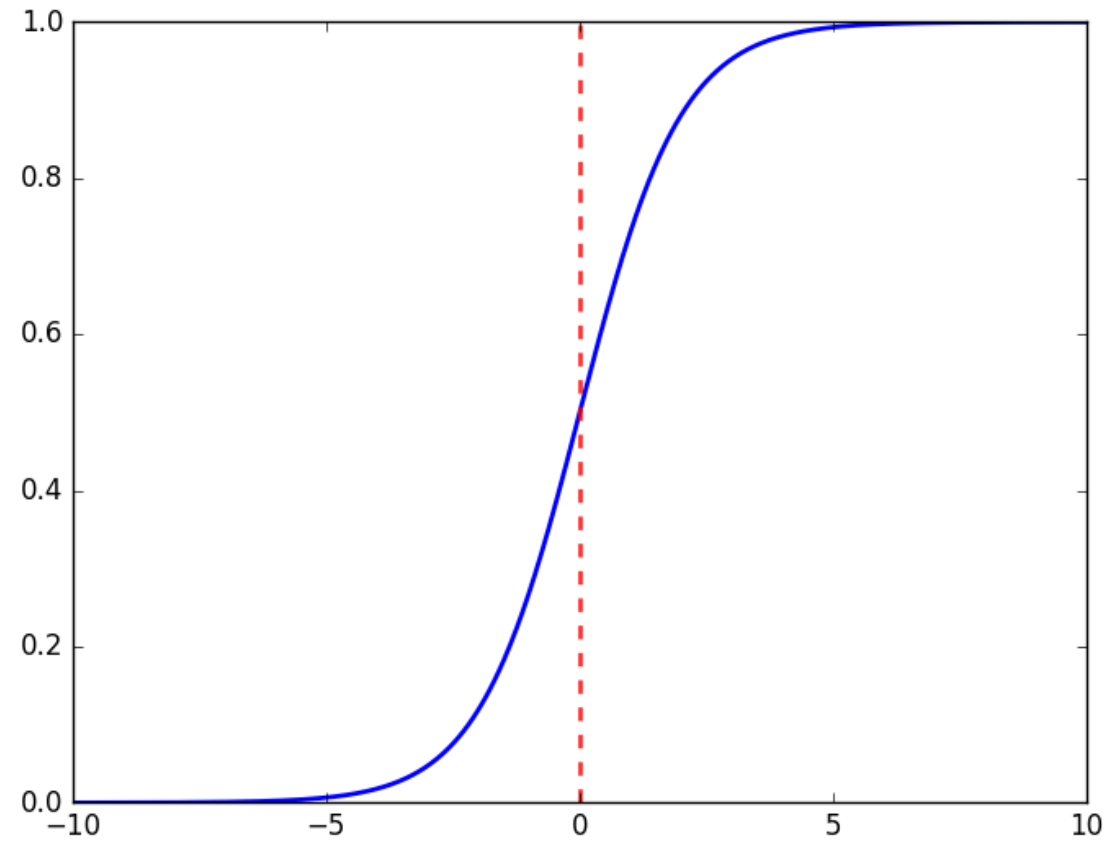
class Net(nn.Module):
    def __init__(self, input_size, hidden_size):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.fc2 = nn.Linear(hidden_size, 1)

    def forward(self, x):
        x = torch.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```

# Binary classification

To obtain the required binary classification, we pass the output  $f(x)$  through another function :

$$g(z) = 1/(1 + e^{-z})$$



The graph of the sigmoid function

# Binary classification

To obtain the required binary classification, we pass the output  $f(x)$  through another function :

$$g(z) = 1/(1 + e^{-z})$$

This function returns an output between 0 and 1. The binary classification is set as follows :

If (  $g(z) \geq 0.5$  ) assign the input the positive class

Else assign the input to the negative class

# Binary classification

To obtain the required binary classification, we pass the output  $f(x)$  through another function :

$$g(z) = 1/(1 + e^{-z})$$

This function returns an output between 0 and 1. The binary classification is set as follows :

If (  $g(z) \geq 0.5$  ) assign the input the positive class

Else assign the input to the negative class

But what do we do in the multi-class classification ?

# Multi-class classification : the softmax function

In the case of multi-class classification, we use the softmax activation function.  
Suppose that we have k classes then the softmax activation function is define by :

$$\text{softmax}(z)_i = \frac{\exp(z_i)}{\sum_{l=1}^k \exp(z_l)}$$

Here  $z_i$  represents the  $i$ th element of the input to softmax, which corresponds to class  $i$ .

# Multi-class classification : the softmax function

In the case of multi-class classification, we use the softmax activation function.  
Suppose that we have  $k$  classes then the softmax activation function is defined by :

$$\text{softmax}(z)_i = \frac{\exp(z_i)}{\sum_{l=1}^k \exp(z_l)}$$

Here  $z_i$  represents the  $i$ th element of the input to softmax, which corresponds to class  $i$ .  
The result is a vector containing the probabilities that sample  $x$  belong to each class.

# Multi-class classification : the softmax function

In the case of multi-class classification, we use the softmax activation function.  
Suppose that we have  $k$  classes then the softmax activation function is defined by :

$$\text{softmax}(z)_i = \frac{\exp(z_i)}{\sum_{l=1}^k \exp(z_l)}$$

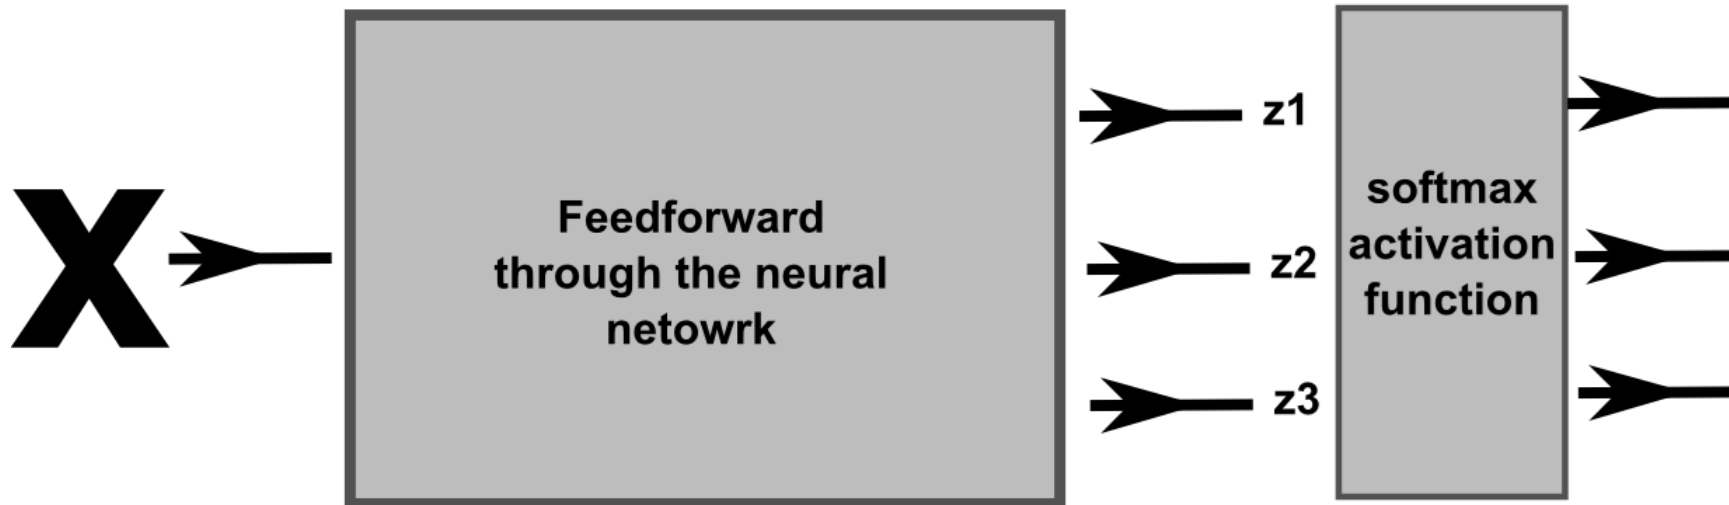
Here  $z_i$  represents the  $i$ th element of the input to softmax, which corresponds to class  $i$ .  
The result is a vector containing the probabilities that sample  $x$  belong to each class.  
The output is the class with the highest probability.

# Multi-class classification : the softmax function

In the case of multi-class classification, we use the softmax activation function.  
Suppose that we have k classes then the softmax activation function is define by :

$$\text{softmax}(z)_i = \frac{\exp(z_i)}{\sum_{l=1}^k \exp(z_l)}$$

Here  $z_i$  represents the  $i$ th element of the input to softmax, which corresponds to class  $i$ .  
The result is a vector containing the probabilities that sample  $x$  belong to each class.  
The output is the class with the highest probability.





# The softmax function in Python

The softmax function is a mathematical function used to convert a vector of real numbers into a probability distribution.

It takes an input vector and returns another vector of the same length, where each element is transformed to a value between 0 and 1, representing the probability of that element being selected. In simple terms, the softmax function normalizes the input vector and makes it easier to interpret as probabilities. Here's a Python example:

```
import numpy as np

def softmax(x):
    exp_values = np.exp(x)
    probabilities = exp_values / np.sum(exp_values)
    return probabilities

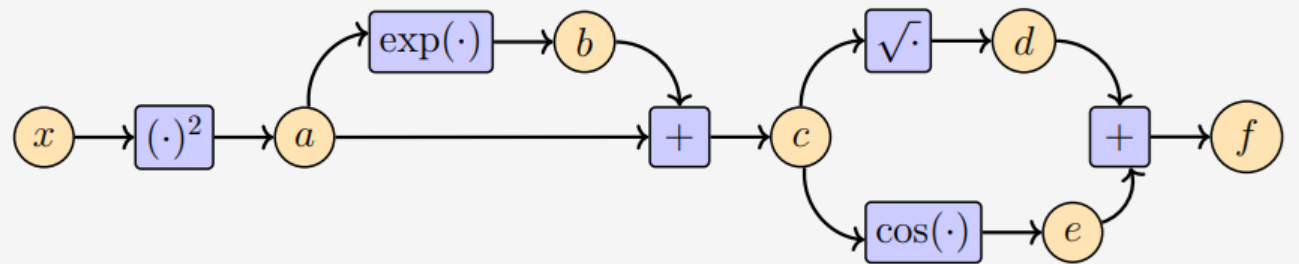
input_vector = np.array([2.0, 1.0, 0.5])
output_vector = softmax(input_vector)
print(output_vector)

[0.62842832 0.2312239 0.14034778]
```

# What is a computational graph ?

$$f(x) = \sqrt{x^2 + \exp(x^2)} + \cos(x^2 + \exp(x^2))$$

$$\begin{aligned} a &= x^2, \\ b &= \exp(a), \\ c &= a + b, \\ d &= \sqrt{c}, \\ e &= \cos(c), \\ f &= d + e. \end{aligned}$$



Computation graph of  $f$

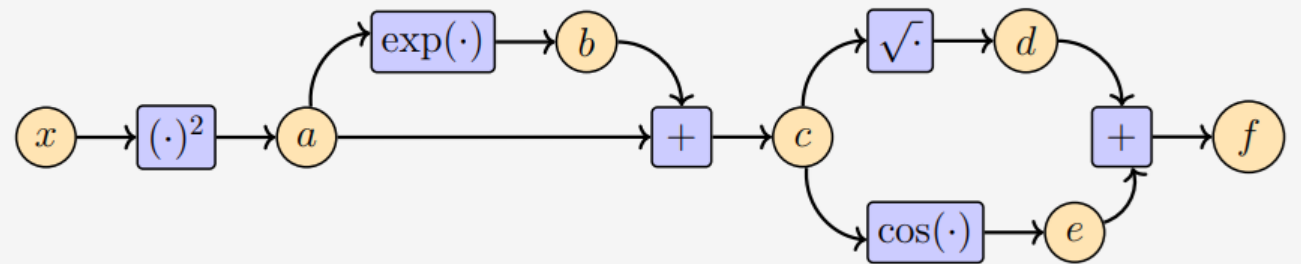
Image source

# What is a computational graph ?

- A computational graph, also known as a computational or **directed acyclic graph** (DAG), is a directed graph that represents a computational process or a sequence of computations.
- It is a graph structure where nodes represent operations or computations, and directed edges represent dependencies between these operations.
- Note: yellow nodes in the graph here are placeholders and not really part of the computational graph. They get executed when we insert a certain input to the graph

$$f(x) = \sqrt{x^2 + \exp(x^2)} + \cos(x^2 + \exp(x^2))$$

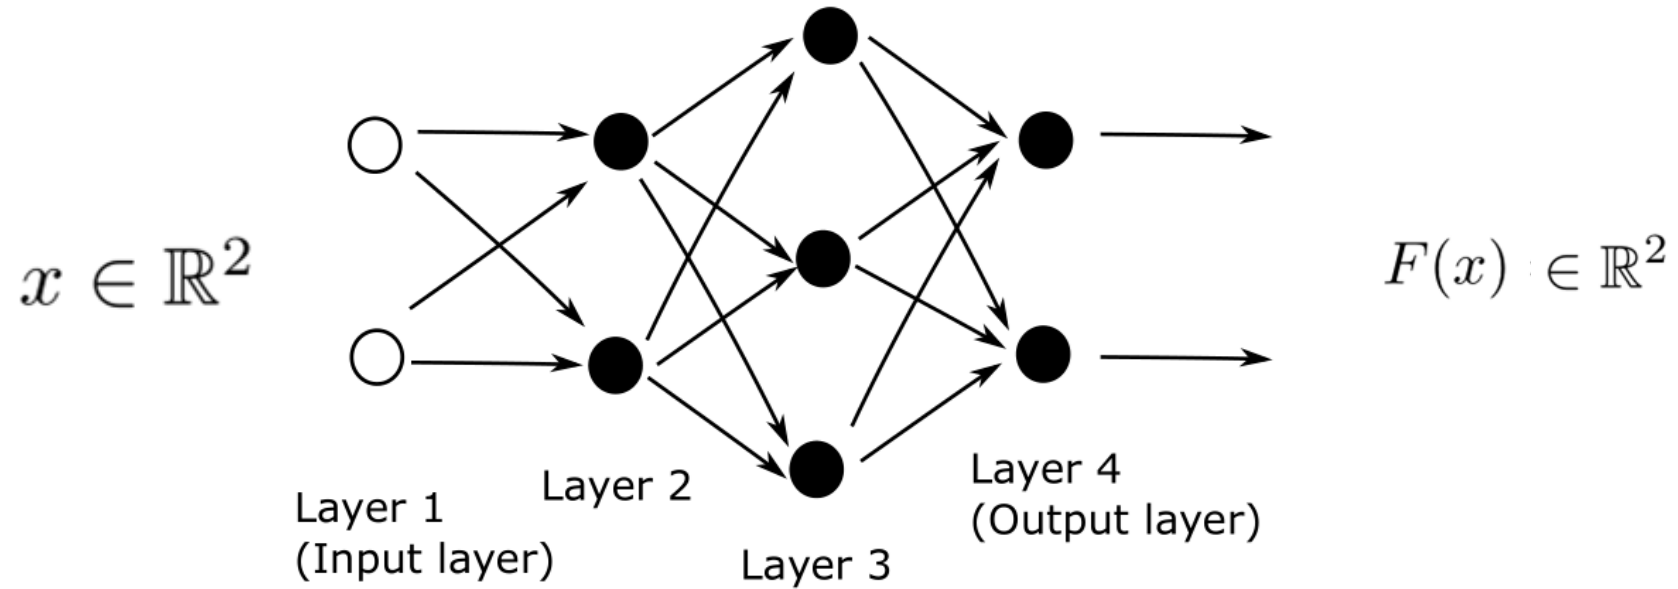
$$\begin{aligned} a &= x^2, \\ b &= \exp(a), \\ c &= a + b, \\ d &= \sqrt{c}, \\ e &= \cos(c), \\ f &= d + e. \end{aligned}$$



Computation graph of  $f$

Image source

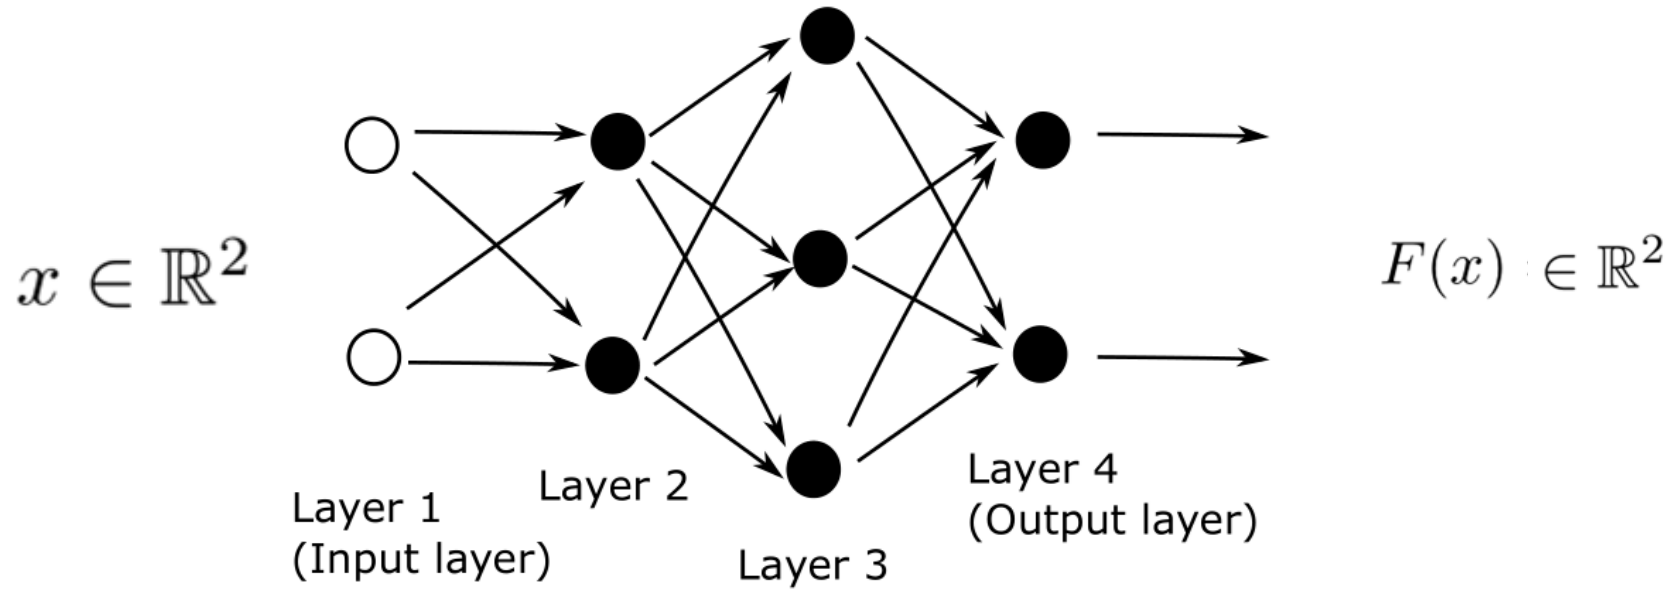
# Neural Networks are computational graphs



Neural networks can be considered as computational graphs.

Why this is a useful fact? Modern DL packages such as tensorflow and pytorch use this fact for automatic differentiation.

# Neural Networks are computational graphs


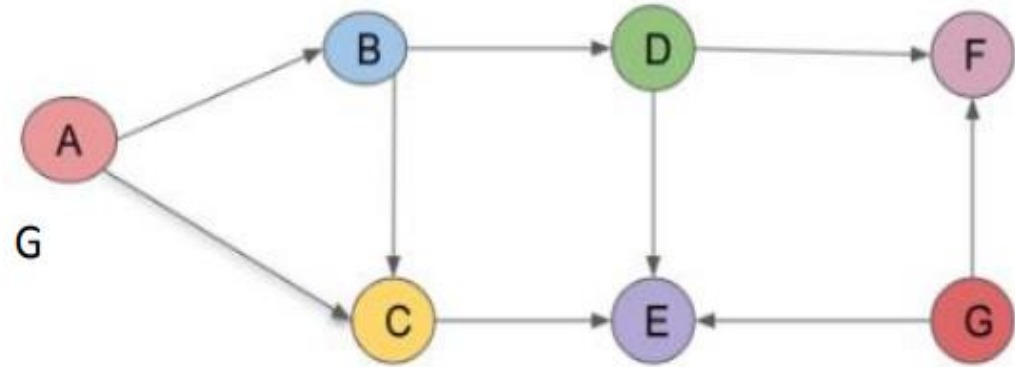


**Key fact** : feedforward computation of a NN is defined to be the computations that one executes on a computational graph that defines that network given an input and a topological order of the nodes of the computational graph of a NN

## Recall topological sort

Recall that a topological sort of a DAG is a linear ordering of its vertices such that for every directed edge  $uv$  from vertex  $u$  to vertex  $v$ ,  $u$  comes before  $v$  in the ordering.

Example



A topological sort of the graph G

## Backprop in nutshell : Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 2, y = 1, z = 0$$

Forward prop step

$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$


Back prop step (local gradients)

$$\partial a / \partial x = 1, \partial a / \partial y = 1$$

$$\partial b / \partial y = \mathbf{1}(y > z), \partial b / \partial z = \mathbf{1}(z > y) = 0$$

$$\partial f / \partial a = b = 3, \\ \partial f / \partial b = a = 1$$

$$\partial f / \partial f = 1$$



# Backprop in nutshell : Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 2, y = 1, z = 0$$

Forward prop step

$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$


Back prop step (local gradients)

$$\partial a / \partial x = 1, \partial a / \partial y = 1$$

$$\partial b / \partial y = \mathbf{1}(y > z), \partial b / \partial z = \mathbf{1}(z > y) = 0$$

$$\partial f / \partial a = b = 3, \\ \partial f / \partial b = a = 1$$

$$\partial f / \partial f = 1$$





# Backprop in nutshell : Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 2, y = 1, z = 0$$

Forward prop step

$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$


Back prop step (local gradients)

$$\partial a / \partial x = 1, \partial a / \partial y = 1$$

$$\partial b / \partial y = \mathbf{1}(y > z), \partial b / \partial z = \mathbf{1}(z > y) = 0$$

$$\partial f / \partial a = b = 3, \\ \partial f / \partial b = a = 1$$

$$\partial f / \partial f = 1$$



# Backprop in nutshell : Example

$$f(x, y, z) = (x + y) \max(y, z)$$

$$x = 2, y = 1, z = 0$$

Forward prop step

$$a = x + y$$

$$b = \max(y, z)$$

$$f = ab$$


Back prop step (local gradients)

$$\partial a / \partial x = 1, \partial a / \partial y = 1$$

$$\partial b / \partial y = \mathbf{1}(y > z), \partial b / \partial z = \mathbf{1}(z > y) = 0$$

$$\partial f / \partial a = b = 3,$$
$$\partial f / \partial b = a = 1$$

$$\partial f / \partial f = 1$$



# Backprop in nutshell : Example

$f(x, y, z) = (x + y) \max(y, z)$   
 $x = 2, y = 1, z = 0$

Forward prop step

$a = x + y$   
 $b = \max(y, z)$   
 $f = ab$


Back prop step (local gradients)

$\partial a / \partial x = 1, \partial a / \partial y = 1$

$\partial b / \partial y = \mathbf{1}(y > z), \partial b / \partial z = \mathbf{1}(z > y) = 0$

$\partial f / \partial a = b = 3,$   
 $\partial f / \partial b = a = 1$

$\partial f / \partial f = 1$



# Backprop in nutshell : Example

$f(x, y, z) = (x + y) \max(y, z)$   
 $x = 2, y = 1, z = 0$

Forward prop step

$a = x + y$   
 $b = \max(y, z)$   
 $f = ab$


Back prop step (local gradients)

$\partial a / \partial x = 1, \partial a / \partial y = 1$


$\partial b / \partial y = \mathbf{1}(y > z), \partial b / \partial z = \mathbf{1}(z > y) = 0$

$\partial f / \partial a = b = 3,$   
 $\partial f / \partial b = a = 1$

$\partial f / \partial f = 1$



# Backprop in nutshell



# Backprop in nutshell

More general :

## Automatic Differentiation – Reverse Mode (aka. Backpropagation)

### Forward Computation

1. Write an **algorithm** for evaluating the function  $y = f(\mathbf{x})$ . The algorithm defines a **directed acyclic graph**, where each variable is a node (i.e. the “**computation graph**”)
2. Visit each node in **topological order**.  
For variable  $u_i$  with inputs  $v_1, \dots, v_N$ 
  - a. Compute  $u_i = g_i(v_1, \dots, v_N)$
  - b. Store the result at the node

### Backward Computation

1. **Initialize** all partial derivatives  $dy/du_i$  to 0 and  $dy/dy = 1$ .
2. Visit each node in **reverse topological order**.  
For variable  $u_i = g_i(v_1, \dots, v_N)$ 
  - a. We already know  $dy/du_i$
  - b. Increment  $dy/dv_j$  by  $(dy/du_i)(du_i/dv_j)$   
(Choice of algorithm ensures computing  $(du_i/dv_j)$  is easy)

**Return** partial derivatives  $dy/du_i$  for all variables

The approximation power of neural networks

# Approximation Theorems using shallow NN

Let  $\varphi : \mathbb{R} \rightarrow \mathbb{R}$  be a constant, bounded, and continuous function.

(You may think about this function as Relu if you like). Consider summation of the form :

$$\sum_{i=1}^N v_i \varphi (w_i x + b_i)$$



Real numbers (the weights)




# Approximation Theorems using shallow NN

Let  $\varphi : \mathbb{R} \rightarrow \mathbb{R}$  be a constant, bounded, and continuous function.

(You may think about this function as Relu if you like). Consider summation of the form :

$$\sum_{i=1}^N v_i \varphi(w_i x + b_i)$$

Real numbers (the weights)



Let  $\varphi(x)=1$  when  $x \geq 0$  and zero otherwise and consider:  $0.5(\varphi(x-1) + \varphi(-x+3))$


# Approximation Theorems using shallow NN

Let  $\varphi : \mathbb{R} \rightarrow \mathbb{R}$  be a constant, bounded, and continuous function.

Consider summation of the form :

$$\sum_{i=1}^N v_i \varphi(w_i x + b_i)$$

Real numbers (the weights)




$$0.5(\varphi(x-1) + \varphi(-x+3))$$


Let  $\varphi(x)=1$  when  $x \geq 0$  and zero otherwise and consider:

Can you imagine building more complex functions if we have more Summations and maybe vary the weights ? –what are the functions you can build?

# Approximation Theorems using shallow NN




[https://en.wikipedia.org/wiki/Riemann\\_sum](https://en.wikipedia.org/wiki/Riemann_sum)



Familiar ? Riemann sum from Calc?

# Approximation Theorems using shallow NN

Question : can you imagine building more complex functions if we have more summations and maybe vary the weights  $w_i$ 's ? what are the functions you can build?




$$\sum_{i=1}^N v_i \varphi(w_i x + b_i)$$

Given a function  $f$  as above, can we find  $v_i$ 's,  $w_i$ 's, and  $b_i$ 's such that the summation above is as close we like to  $f$  ?

This is the essence of the universal approximation theorem : it can always be done.

# Approximation Theorems using shallow NN

Question : can you imagine building more complex functions if we have more summations and maybe vary the weights  $w_i$ 's ? what are the functions you can build?



$$\sum_{i=1}^N v_i \varphi(w_i x + b_i)$$

Given a function  $f$  as above, can we find  $v_i$ 's,  $w_i$ 's, and  $b_i$ 's such that the summation above is as close we like to  $f$  ?

It turns out that the answer is yes as long as we are willing to increase  $N$  (increase number of kernels).  
Let's [see](#) a few examples.

## Approximation Theorems using shallow NN

It turns out that this theorem generalizes to higher dimension the same way. More precisely, the following summations:

$$\sum_{i=1}^N v_i \varphi(w_i^T x + b_i) \quad \begin{array}{l} w_i \in \mathbb{R}^m \\ v_i, b_i \in \mathbb{R} \end{array}$$

can approximate any continuous real valued function on  $[0,1]^m$ .

# Approximation Theorems using shallow NN

It turns out that this theorem generalizes to higher dimension the same way. More precisely, the following summations:

$$\sum_{i=1}^N v_i \varphi(w_i^T x + b_i) \quad \begin{array}{l} w_i \in \mathbb{R}^m \\ v_i, b_i \in \mathbb{R} \end{array}$$

can approximate any continuous real valued function on  $[0,1]^m$ .

**Universal approximation theorem.** Let  $\varphi : \mathbb{R} \rightarrow \mathbb{R}$  be a nonconstant, **bounded**, and **continuous** function (called the *activation function*). Let  $I_m$  denote the  $m$ -dimensional **unit hypercube**  $[0, 1]^m$ . The space of real-valued continuous functions on  $I_m$  is denoted by  $C(I_m)$ . Then, given any  $\varepsilon > 0$  and any function  $f \in C(I_m)$ , there exist an integer  $N$ , real constants  $v_i, b_i \in \mathbb{R}$  and real vectors  $w_i \in \mathbb{R}^m$  for  $i = 1, \dots, N$ , such that we may define:

$$F(x) = \sum_{i=1}^N v_i \varphi(w_i^T x + b_i)$$

as an approximate realization of the function  $f$ ; that is,

$$|F(x) - f(x)| < \varepsilon$$